

AI NEWS DAILY

Daily Intelligence Digest — 2026-06-25

Total: 17 | AI Trends: 11 | Tech: 4 | Thailand: 0 | Global: 2 | Breakthroughs: 4 | High-urgency: 6

[HIGH] Claude Tag AI Slack

Anthropic 65% @Claude in Slack benchmark AI adoption engineering persistent memory background task execu

[HIGH] OpenAI and Broadcom unveil LLM-optimized inference chip 'Jalapeño'

Custom silicon for inference (not training) reflects the industry's maturing understanding that inference efficiency — not just training capability — is the primary cost driver at scale. The 9-month d

[HIGH] Qualcomm to Acquire Modular — \$3.9B deal to accelerate AI software for edge-to-cloud

This acquisition represents Qualcomm's most significant software bet: Modular's 'write once, run anywhere' approach for AI models is the AI equivalent of what LLVM did for compilers. Lattner's pedigree

[HIGH] AI researchers continue to leave Google for its rivals

The concentration of departures in June 2026 — immediately preceding Anthropic's IPO — suggests pre-IPO equity grants are the primary driver. Google's structure as a large public company makes it stru

AI TRENDS (11 articles)

[HIGH] Claude Tag AI Slack

Anthropic has launched Claude Tag on Slack, a major upgrade to its previous Claude for Slack integration, targeting Enterprise and Team plan subscribers and powered by the Claude Opus 4.8 model. Key capabilities include persistent memory across sessions, cross-user conversation context retention, background task execution that does not require the user to stay active in a thread, and deep integration with third-party tools. Anthropic disclosed a striking internal benchmark: 65% of all code written at the company is now authored through @Claude within Slack, illustrating the depth of AI adoption achievable in engineering workflows. Claude Tag marks a meaningful shift from AI-as-chatbot to AI-as-team-member, operating inside real communication environments rather than isolated chat interfaces. (Source: Blognone)

Source: Blognone | #Anthropic #Claude Tag #Claude Opus 4 #Slack

[HIGH] [BREAKTHROUGH] OpenAI and Broadcom unveil LLM-optimized inference chip 'Jalapeño'

OpenAI and Broadcom have unveiled 'Jalapeño,' OpenAI's first custom AI inference accelerator, developed in just nine months — the fastest ASIC development timeline for high-performance compute on record. Designed exclusively for large language model inference (not general-purpose compute), engineering samples are already running GPT-5.3-Codex-Spark at production target performance. The chip delivers substantially better performance-per-watt than current state-of-the-art alternatives, positioning OpenAI to dramatically reduce the cost of serving its models at scale. OpenAI plans to deploy Jalapeño at gigawatt scale by end-2026 in partnership with Microsoft and other infrastructure partners. Celestica is named as a manufacturing partner alongside Broadcom. This is part of a multi-generation chip roadmap, signaling OpenAI's shift toward full-stack infrastructure control. Greg Brockman framed the move as a 'full-stack infrastructure strategy to make compute more abundant' for AI development globally.

Source: OpenAI Blog | #OpenAI #Broadcom #Jalapeno #AI chip

[HIGH] [BREAKTHROUGH] Qualcomm to Acquire Modular — \$3.9B deal to accelerate AI software for edge-to-cloud

Qualcomm announced a definitive agreement to acquire Modular Inc., the AI software infrastructure company behind the Mojo programming language and the MAX AI platform, for approximately \$3.9 billion. Founded by Chris Lattner — creator of LLVM, Clang, MLIR, and the Swift programming language — Modular built a platform enabling AI models to run efficiently across heterogeneous hardware (CPU, GPU, NPU, custom ASICs) without per-accelerator rewrites. For Qualcomm, the acquisition strengthens its software foundation for generative and agentic AI across data center and edge environments, deepening its data center strategy alongside its established edge silicon leadership. Qualcomm CEO Cristiano Amon called it 'a pivotal moment' as the industry moves toward disaggregated, multi-vendor architectures that demand open, modern software foundations. The deal is expected to close in the second half of 2026 subject to regulatory approvals.

Source: Modular Blog | #Qualcomm #Modular #Mojo #MAX platform

[HIGH] AI researchers continue to leave Google for its rivals

Top Google AI researchers Jonas Adler and Alexander Pritzel are leaving to join Anthropic, deepening a talent exodus from Google DeepMind that has accelerated dramatically in June 2026. Adler and Pritzel played key roles in developing Google's Gemini model. This follows the departure of legendary researcher Noam Shazeer — at Google since 2000, briefly at Character.AI — to OpenAI last week, and Nobel Prize-winning DeepMind director John Jumper's move to Anthropic just days earlier. Analysts note that the upcoming IPOs of both OpenAI and Anthropic are creating powerful equity-based recruiting incentives, making it an unusually attractive time for researchers to switch companies. The pattern raises systemic questions about Google's ability to retain frontier AI talent as its rivals prepare for public listings that could yield enormous equity windfalls for early employees.

Source: TechCrunch | #Google #Anthropic #OpenAI #AI talent

[HIGH] Anthropic accuses Alibaba of campaign to 'brazenly' and 'illicitly' extract Claude AI capabilities

Anthropic sent a formal letter to the U.S. Senate Committee on Banking, Housing, and Urban Affairs accusing Alibaba's AI lab Qwen of orchestrating the largest known model distillation attack on Claude to date. Between April 22 and June 5, 2026, approximately 25,000 fraudulent accounts linked to Alibaba-affiliated operators conducted 28.8 million exchanges with Claude, specifically targeting its most commercially valuable capabilities: software engineering and advanced agentic reasoning. Anthropic called the campaign 'brazen' and 'illicit,' noting it occurred after the White House had publicly warned about distillation attacks as a national security threat. Alibaba's ADR shares fell more than 3% on the news. The letter urges Congress to strengthen antitrust information-sharing rules, tighten AI chip export controls, and impose sanctions on companies conducting distillation attacks. Senators Bill Hagerty and Andy Kim are reportedly preparing legislation in response.

Source: CNBC | #Anthropic #Alibaba #Qwen #distillation

[LOW] Facebook rolls out an AI companion app for creators

Facebook announced it is reimagining its Creator Studio tool as a standalone AI companion app designed to help creators grow their audiences on the platform. Currently in testing with select creators, the app features an AI assistant that provides personalized recommendations based on content style, performance metrics, audience engagement, and creator goals. Creators can ask conversational questions like 'When should I post?' or 'What are people saying in my comments?' and get instant contextual answers without navigating complex dashboards. The app also includes an AI-powered comment tool that surfaces important responses and drafts replies in the creator's own voice. Meta CEO Mark Zuckerberg has said AI-driven efficiencies are enabling the company to build more apps than historically possible, with recent launches including Forum (a Reddit-like community app) and Instants (a disappearing-photos Instagram app). The Creator Studio companion follows the previously reported 'Arena' prediction market app still in development.

Source: TechCrunch | #Meta #Facebook #creator economy #AI assistant

[MEDIUM] [BREAKTHROUGH] MindZero: Learning Online Mental Reasoning With Zero Annotations

MindZero is a new framework accepted to ICML 2026 for teaching AI agents to reason about multi-step tasks without any human annotations. Traditional approaches for training reasoning agents require extensive labeled data

describing thought processes — an expensive and brittle bottleneck. MindZero instead learns to reason purely from interaction signals using a multi-agent setup: one model generates reasoning traces while another evaluates their quality through environmental feedback. The system showed strong generalization across diverse tasks requiring multi-step planning and inference about other agents' intentions. By eliminating annotation dependency, MindZero opens a path toward scalable training pipelines for real-world agent deployments where labeling costs grow proportionally with task complexity. The work is particularly relevant for agentic AI systems that must reason about dynamic, partially observable environments — matching the trajectory of production AI agent deployments. Code and model weights available at scai.cs.jhu.edu/MindZero.

Source: [arXiv](#) | [#arXiv](#) [#ICML 2026](#) [#reasoning](#) [#multi-agent](#)

[MEDIUM] The Deterministic Horizon: When Extended Reasoning Fails and Tool Delegation Becomes Necessary

Accepted to ICML 2026, this paper formalizes a key limitation of extended reasoning models: there exists a class of problems — deterministic problems with a unique correct answer — where additional reasoning steps provide no benefit and the only path to improvement is delegating to external tools. The authors introduce the 'deterministic horizon,' the point beyond which a reasoning LLM's chain-of-thought becomes circular or repetitive rather than productive. Empirical results across mathematics, code execution, and database lookup tasks show models waste up to 60% of extended reasoning tokens on post-horizon steps that contribute nothing to final accuracy. The paper proposes a detection mechanism to identify when the horizon is reached and automatically invoke tool use, significantly improving accuracy and reducing compute waste. This work has direct implications for the design of production agentic systems where reasoning and tool-use are mixed.

Source: [arXiv](#) | [#arXiv](#) [#ICML 2026](#) [#reasoning](#) [#chain-of-thought](#)

[MEDIUM] [BREAKTHROUGH] Threshold-Based Exclusive Batching for LLM Inference

Threshold-Based Exclusive Batching (TEB), accepted to ICML 2026, is a novel inference scheduling algorithm that addresses performance degradation in LLM serving caused by heterogeneous request lengths within a batch. Traditional continuous batching treats all requests equally, causing short requests to wait behind long ones and degrading throughput under mixed-workload conditions. TEB dynamically separates requests into exclusive batch groups based on an adaptive length threshold, preventing cross-contamination of latency profiles. Experiments on production-scale LLM serving infrastructure show TEB achieves up to 2.3× throughput improvement over continuous batching while maintaining sub-100ms time-to-first-token at the 95th percentile. The approach is model-agnostic and compatible with existing KV-cache management and speculative decoding techniques, making it directly deployable in current production stacks without architectural changes to the underlying model.

Source: [arXiv](#) | [#arXiv](#) [#ICML 2026](#) [#LLM inference](#) [#batching](#)

[HIGH] Hidden Thoughts Are Not Secret: Reasoning Trace Exposure in LLMs

Researchers from National Taiwan University and UC Berkeley reveal that the internal reasoning traces of chain-of-thought LLMs are not meaningfully private and can be reverse-engineered by a sufficiently motivated adversary. Even when a model's extended reasoning output is hidden from the user — a common deployment pattern to protect proprietary prompts or confidential data — an adversary with black-box API access can reconstruct the reasoning trace with high fidelity by systematically probing the model with carefully designed follow-up queries. This has significant implications for AI systems that use private reasoning chains to handle confidential data or proprietary business logic. The authors propose a cryptographic watermarking scheme for reasoning traces that enables detection of reconstruction attempts. The work highlights a fundamental tension between useful reasoning transparency and reasoning confidentiality in deployed AI systems, particularly relevant for enterprise deployments handling sensitive data.

Source: [arXiv](#) | [#arXiv](#) [#LLM security](#) [#reasoning traces](#) [#chain-of-thought](#)

[MEDIUM] MOSAIC: Modular Orchestration for Structured Agentic Intelligence and Composition

MOSAIC presents a framework for modular orchestration of structured agentic AI systems that enables composable, reusable agent architectures for complex task completion. Unlike monolithic agent designs, MOSAIC treats agent

Source: arXiv | #arXiv #AI agents #multi-agent #orchestration

Source: Blognone | #Meta #prediction market #Polymarket #Arena

Source: Blognone | #Rockstar Games #GTA 6 #GTA VI #gaming

Source: TechCrunch | #Figma #AI design #code layers #animation

Agility Robotics, the humanoid robotics startup spun out of Oregon State University in 2015, announced plans to go public via merger with Churchill Capital Corp XI SPAC, valuing the company at approximately \$2.5 billion and

expecting to generate over \$620 million in proceeds including \$200 million from new institutional investors. Agility is best known for Digit, its bipedal humanoid robot currently deployed across nine commercial customer sites including Toyota Motor Manufacturing Canada, Schaeffler, GXO, and Mercado Libre. The company has secured over \$300 million in multi-year orders for its next-generation Digit v5, with a pipeline of 30+ potential enterprise customers evaluating large-scale deployments. The combined company will trade as AGLT on a North American exchange. Backers include Amazon, Nvidia, and SoftBank Vision Fund 2. CEO Peggy Johnson cited humanoid robots as 'a critical driver of productivity, supply chain resilience, and American technology leadership.'

Source: TechCrunch | #Agility Robotics #humanoid robot #Digit #SPAC

THAILAND AI & TECH (0 articles)

No articles in this category today.

GLOBAL AI & TECH (2 articles)

[MEDIUM] Summer Davos AI

At the World Economic Forum's Summer Davos 2026 in Dalian, China, South Korea's Prime Minister announced a national strategy to become the world's third-largest AI superpower, paired with a Universal Basic Income proposal as a long-term policy response to workforce displacement driven by AI automation. The strategy calls for heavy investment in AI infrastructure including semiconductor chips, data centers, and development of large language models natively trained on Korean. The UBI proposal marks a concrete legislative step—moving the AI displacement conversation from academic debate to formal policy agenda. South Korea's announcement came amid intensifying technology competition between the United States and China, with Seoul positioning itself as a neutral but strategically capable third hub for global AI development. (Source: Techsauce)

Source: Techsauce | #South Korea #AI policy #Summer Davos #WEF

[MEDIUM] ?

A Techsauce analysis published during Summer Davos 2026 examines China's transformation from "world factory" to a global technology innovator, driven by sustained R&D; investment growth, a maturing startup ecosystem, and systematic state-backed support for innovation. China now competes at the leading edge in artificial intelligence, robotics, clean energy, and electric vehicles—no longer operating as a fast follower of Western technology. Hosting Summer Davos 2026 further underscores China's expanding influence in international economic forums. Chinese AI models such as DeepSeek and Qwen are already competitive with US counterparts on benchmarks at a fraction of the cost, directly challenging assumptions about American AI dominance and reshaping competitive dynamics for Southeast Asian technology ecosystems including Thailand. (Source: Techsauce)

Source: Techsauce | #China #innovation #AI #geopolitics