

AI NEWS DAILY — 2026-06-26

25 articles | AI News Pipeline v2.0

TOTAL: 23

AI: 12

TECH: 8

THAI: 1

GLOBAL: 2

AI TRENDS (12)

Anthropic Claude Tag @Claude
Slack

TechTalkThai | Thai | HIGH | #Anthropic #Claude #Slack #AI Agent

Anthropic Claude Tag AI
Slack @Claude
channel
multiplayer
ambient mode Claude
Anthropic 65% Claude Tag
Claude Tag Opus 4.8 RBAC
Claude Enterprise Team
beta Claude in Slack
30

Claude Tag AI one-on-one chat multiplayer team collaboration
adoption AI

OpenAI Broadcom Jalapeño Inference LLM

TechTalkThai | Thai | MED | #OpenAI #Broadcom #Jalapeño #LLM Chip

OpenAI Broadcom Jalapeño
Inference LLM
NVIDIA GPU
supply Jalapeño inference
general-purpose AI accelerators serve LLM models
production throughput latency
custom silicon AI
hardware performance
Jalapeño chip custom AI silicon general-purpose GPU
LLM inference NVIDIA AI infrastructure

Deep Tech AI

Techsauce | Thai | LOW | #Deep Tech #AI Infrastructure #Economic Infrastructure #WEF

'Deep Tech, Shallow Markets' Summer Davos 2026
Deep Tech
AI
Deep Tech Biotech, Quantum Computing, Advanced Materials commercialize
AI cycle position economic foundation
Deep Tech economic infrastructure
funding landscape AI startup

How agents are transforming work

OpenAI | International | HIGH | #OpenAI #Codex #AI Agents #Agentic AI

OpenAI published a new Economic Research paper documenting how agentic AI — specifically Codex — has fundamentally shifted work patterns. By May 2026, 80.6% of Codex individual users made at least one request corresponding to over 30 minutes of human work, with 25.6% making requests equivalent to more than 8 hours. Codex now accounts for 99.8% of weekly output tokens generated internally at OpenAI, with non-technical departments like Legal and Recruiting adopting it as their primary AI tool around April 2026. Non-developer adoption grew 137x for individual users and 189x for organizational users since August 2025. The paper documents how users at the 99th percentile regularly generate over 60 hours of Codex agent turns per day across parallel agents. This marks a structural shift from single chatbot interactions to delegated, long-horizon agentic work across every function.

This paper is the strongest evidence yet that agentic AI has crossed a critical adoption threshold — even non-technical workers are delegating multi-hour tasks autonomously. This will accelerate enterprise restructuring around AI agents within 12 months.

Patronus AI lands \$50M to build 'digital worlds' that stress-test AI agents

TechCrunch | International | HIGH | #Patronus AI #AI Agents #Agent Testing #Series B

Patronus AI, founded by former Meta AI researchers Anand Kannappan and Rebecca Qian, raised a \$50M Series B led by Greenfield Partners with participation from Notable Capital, Lightspeed, Datadog, and Samsung — bringing total funding to \$70M. The company builds 'digital world models,' simulated replicas of websites and internal enterprise systems, where AI agents are stress-tested using reinforcement learning. Virtually every frontier AI lab is now a customer, and revenue has grown 15-fold over the past year amid what investors call 'nearly insatiable' demand. Patronus' approach parallels how Waymo trained autonomous vehicles in synthetic worlds before real deployment — but for AI agents prone to finding 'shortcuts' that complete tasks incorrectly. Current focus areas are software engineering and finance use cases, with ambitions to expand to non-verifiable task categories. Competing primarily against internal evaluation teams at AI labs, not other startups.

Agent evaluation is becoming critical infrastructure as enterprises deploy AI agents. Patronus represents a new category — AI quality assurance — that will be mandatory before any autonomous agent can be trusted with financial, legal, or operational tasks.

Anthropic's Claude is winning over paid consumers, a market owned by ChatGPT

TechCrunch | International | ■ MED | #Anthropic #Claude #OpenAI #ChatGPT

Data from Indagari analyzing billions of anonymized credit card transactions from 28 million US consumers shows Anthropic's Claude growing paying subscribers approximately 75% since January 2026 — with growth continuing even after a spike in March when Anthropic refused to allow model use for mass surveillance. On DataCamp's online AI education platform with 20M users, 'Claude' is now the most searched term, surpassing even the term 'AI' itself. Claude course demand on DataCamp has grown 18x in the last 30 days, outpacing ChatGPT demand 3-to-1 among self-directed learners. Both Anthropic and OpenAI are preparing for IPOs. Despite the impressive growth trajectory, ChatGPT remains the dominant AI assistant by a wide margin in absolute numbers — the data shows market share dynamics are shifting but ChatGPT's lead is structural.

■ *Claude's consumer momentum is a strategic signal for the broader AI market — the 'responsible AI' positioning combined with the Fable 5 controversy is generating 'forbidden fruit' brand affinity that ChatGPT cannot replicate, reshaping the competitive landscape ahead of two potential AI IPOs.*

General Intuition's \$2.3B bet that video games can train AI agents for the real world

■ TechCrunch | International | ■ MED | #General Intuition #AI Agents #Video Games #AI Training

General Intuition is pursuing a \$2.3B valuation thesis that video game environments represent the optimal training ground for AI agents that can operate in the real world. The company's approach leverages the rich, interactive, physically accurate, and endlessly varied nature of video games as a proxy for real-world complexity — far cheaper and more scalable than real-world data collection. The startup raised \$300M in a recent round and is betting that the gap between simulated and real-world environments for AI agents is closing rapidly, similar to how reinforcement learning in games led to breakthrough results in robotics and strategy AI. If proven correct, video game-trained agents could dramatically accelerate deployment timelines for autonomous agents across industries. The company is competing in a crowded simulation-for-AI-training space but differentiates through the scale and diversity of game environments.

■ *Video game-as-AI-training-ground is a high-leverage bet — if it works at scale, it could 10x the speed of capable agent development by replacing expensive real-world data collection with infinite synthetic diversity.*

DeepSeek-V4: Towards Highly Efficient Million-Token Context Intelligence

■ ArXiv | Research | ■ HIGH | #DeepSeek #DeepSeek-V4 #MoE #Long Context

DeepSeek-AI introduces DeepSeek-V4, comprising two Mixture-of-Experts models: DeepSeek-V4-Pro (1.6T parameters total, 49B activated) and DeepSeek-V4-Flash (284B parameters, 13B activated). Both natively support one million token context windows. Key architectural innovations include Compressed Sparse Attention (CSA) and Heavily Compressed Attention (HCA) for long-context efficiency, Manifold-Constrained Hyper-Connections (mHC) as improved residual connections, and the Muon optimizer for training stability and faster convergence. Both models are pre-trained on more than 32 trillion tokens. At 1M token context, DeepSeek-V4-Pro requires only 27% of single-token inference FLOPs and 10% of KV cache compared to DeepSeek-V3.2 — enabling routine 1M-context support. DeepSeek-V4-Pro-Max (maximum reasoning effort) 'redefines state-of-the-art for open models.' Model checkpoints are publicly released on HuggingFace.

■ *DeepSeek-V4 is the most significant open model release of 2026 to date — a 1.6T-parameter open-weight MoE with 1M token context that reportedly beats frontier closed models while being freely deployable. This shifts the open vs. closed model competitive dynamic dramatically and gives every enterprise access to enterprise-grade long-context reasoning without API dependency.*

Reliability without Validity: A Systematic, Large-Scale Evaluation of LLM-as-a-Judge Models Across Agreement, Consistency, and Bias

ArXiv | Research | ■ HIGH | #LLM-as-a-Judge #AI Evaluation #Bias #Position Bias

Researchers conduct the largest systematic evaluation of LLM-as-a-Judge to date: 21 judges from 9 providers (including frontier models through April 2026) evaluated across MT-Bench, JudgeBench, and RewardBench using 541,000 individual judgments over 118 runs. Four critical findings: (1) exact-match agreement overstates discriminative ability by 33-41 percentage points versus Cohen's kappa, meaning most published LLM evaluation benchmarks are inflated; (2) judge rankings shift by up to 14 positions across benchmarks, making leaderboard comparisons meaningless; (3) two production-deployed judges simultaneously show high test-retest reliability (>0.95) and severe position bias (>0.10) — a 'consistency-bias paradox' where judges consistently favor whatever response appears first; and (4) verbosity bias is smaller than feared (<0.011). The paper proposes a Minimum Viable Validation Protocol for future evaluations. This study fundamentally challenges how the industry measures AI quality.

■ *This paper is a landmark challenge to the entire AI evaluation ecosystem — if LLM-as-a-Judge systematically inflates scores and suffers position bias even in production models, every benchmark comparison, RLHF training signal, and product quality claim that relies on LLM judges is suspect. This will force a reckoning in how AI labs report progress.*

AtomMem: Building Simple and Effective Memory System for LLM Agents via Atomic Facts

ArXiv | Research | ■ MED | #AI Agents #Memory #LLM #Long-term Memory

AtomMem introduces a long-term memory system for LLM agents that addresses the fundamental limitation of fixed context windows. The system uses a Fact Executor component to selectively extract high-value 'atomic facts' from long-form interactions, storing them as dense, efficient memory representations rather than raw conversation history. These atomic facts are organized into hierarchical event structures and temporal profiles that capture coherent episodic context and track dynamically evolving user attributes over time. During retrieval, an associative memory graph connects fragmented memories across sessions. AtomMem achieves state-of-the-art performance on the LoCoMo benchmark across reasoning tasks requiring long-horizon memory, offering a scalable and economically viable approach for personalized AI agents. The atomic facts representation is significantly more token-efficient than storing full conversation history, making true long-term memory feasible for production deployment.

■ *AtomMem's atomic facts approach provides a practical solution to one of the core limitations of production AI agents — the inability to remember across sessions at scale. This architecture directly enables the kind of persistent, personalized agent behavior that enterprise deployments require.*

Trustworthy Multi-Agent Systems: Mitigating Semantic Drift with the Argent Signaling Protocol

ArXiv | Research | ■ HIGH | #Multi-Agent #AI Safety #Trust #LLM Agents

The Argent Signaling Protocol (ASP) introduces a compact machine-readable header that accompanies every AI-generated response in multi-agent LLM systems with structured quality signals: certainty (@C), grounding (@G), stochasticity (@S), and an assumption index. These signals let orchestration controllers distinguish 'repairable failures' (wrong answer, right material — retry) from 'containment failures' (ungrounded output — halt) rather than blindly retrying all errors. Evaluation on a 27-question document-grounded QA benchmark using local GGUF models shows dramatic improvements: Qwen 0.8B pass rate rises from 11.1% to 33.3%; grounding coverage improves from 36.7% to 65.4%. In multi-agent mode, an ASP sidecar blocks 100% of ungrounded upstream outputs from reaching downstream agents (24/27 blocked, zero ungrounded propagations). This directly addresses the 'semantic drift' problem in pipelines where ungrounded outputs compound across agent steps.

■ *ASP addresses one of the most dangerous failure modes in production multi-agent systems — hallucinations that compound as they propagate downstream through agent pipelines. A simple protocol that makes quality signals explicit could become critical infrastructure for any enterprise AI deployment.*

ArXiv | Research | ■ LOW | #Test-Time Scaling #Reasoning #Verification #Process Reward Model

■ GRACE answers one of the most practical questions in production LLM deployment: when to pay for expensive step-level process reward models versus cheap outcome reward models. The phase transition result gives a principled decision rule that could save significant compute costs for teams using test-time scaling.

[illegible]

The word cloud contains the following terms:

- Salesforce
- Agentforce Help Agent
- AI
- Pay-per-resolution
- \$2
- AI
- Help.Salesforce.com
- 4.3
- 70%
- Agentforce Trust Layer
- 2026
- Pay-per-resolution game changer
- AI business model
- pricing strategy
- adoption
- ROI

TechTalkThai | Thai | ■ MED | #IBM #FlashSystem #Agentic AI #Enterprise Storage

IBM FlashSystem.ai x600 Agentic AI enterprise storage
hardware upgrade
storage
90% Agentic AI FlashSystem.ai anomaly
ransomware 60
enterprise storage AI hardware
infrastructure add-on layer
Agentic AI enterprise storage hardware paradigm shift
infrastructure management human-driven AI-driven

Token Monetization

Huawei ■■■■■■ Token Monetization ■■■■■ MWC Shanghai 2026 ■■■■■■
 bandwidth ■■■■■■ AI ■■■■■ token ■■■■■■
 ■■■■■■ AI ■■■■■ commodity service ■■■■■■
 5G-A ■■■■■ monetize ■■■■■ monetize data ■■■■■ 4G
 ■■■■■ operator ■■■■■ differentiated services ■■■ revenue stream
 ■■■■■ connectivity ■■■■■ Huawei ■■■■■ pilot ■■■ telecom
 partners ■■■■■

Token Monetization business model telecom adopt operator AIS, True, DTAC

■ TechCrunch | International | ■ MED | #Databricks #AI Power #Energy Efficiency #AI Infrastructure

■ AI energy efficiency is the single most important infrastructure problem of 2026 — a credible 1,000x improvement would make AI economically deployable at scales that are currently impossible, and would shift the competitive landscape entirely away from companies with energy access advantages.

TechCrunch | International | ■ HIGH | #Notion #Email #AI Agents #Product Discontinuation

■ *Notion Mail's shutdown is an early signal of 'agentic disruption' — entire software product categories being made obsolete not by better competitors but by AI agents that eliminate the underlying user task. Email is just the first; document editing, scheduling, and project management are likely next.*

Netris raises \$15M Series A from a16z to help AI neoclouds go live faster

TechCrunch | International | ■ LOW | #Netris #a16z #Neocloud #AI Infrastructure

Netris has raised a \$15M Series A from Andreessen Horowitz (a16z) to build network automation software specifically designed for the new generation of AI-optimized 'neocloud' data centers. As demand for AI compute has exploded, a new class of cloud providers specializing exclusively in GPU infrastructure has emerged — CoreWeave, Lambda Labs, and others — and they need to provision and orchestrate networks far more rapidly than traditional cloud companies. Netris automates the network configuration, provisioning, and management layer that typically requires highly specialized network engineers, enabling neoclouds to bring new GPU clusters online in hours rather than weeks. The a16z backing signals strong institutional conviction in the neocloud sector as a long-term category separate from hyperscalers.

■ *Neocloud infrastructure software is a new category positioned between hyperscalers and bare-metal GPU providers, and a16z's bet on Netris signals it sees neoclouds as a durable structural force in AI infrastructure — not a transitional category.*

Adobe acquires image and video enhancement tool maker Topaz Labs

TechCrunch | International | ■ LOW | #Adobe #Topaz Labs #AI Image Enhancement #AI Video

Adobe has acquired Topaz Labs, the company behind widely-used AI-powered image upscaling and video enhancement tools including Topaz Photo AI, Topaz Video AI, and Sharpen AI. Topaz Labs' technology is known for best-in-class AI upscaling and noise reduction, with a strong professional user base in photography and video production. For Adobe, this acquisition fills a gap in its Creative Suite with AI enhancement capabilities that many users had been accessing through third-party tools. The integration into Photoshop, Lightroom, and Premiere Pro will bring AI-powered 2K-to-8K upscaling, motion blur correction, and noise reduction natively into the Adobe ecosystem. The acquisition reflects Adobe's broader strategy to defend its creative software moat against generative AI competitors by deepening AI capabilities across its product line.

■ *Adobe's acquisition of Topaz Labs signals the consolidation phase of AI creative tools — standalone AI enhancement apps will be absorbed into platform suites, and the competitive moat for best-in-class AI media processing is shifting to whoever holds the largest user base rather than the best algorithm.*

Pruning via Causal Attribution Preserves Reasoning Performance in Large Language Models

ArXiv | Research | ■ LOW | #LLM Pruning #Model Efficiency #Reasoning #Causal Attribution

Causal Attribution Pruning (CAP) is a training-free method for reducing LLM inference cost while preserving reasoning performance. Rather than using magnitude or activation-based pruning criteria, CAP measures each attention head's causal impact on reasoning tasks by masking heads during forward passes on a calibration set and measuring performance degradation. These causal scores guide fine-grained weight pruning of projection matrices. CAP achieves up to 61% relative accuracy gains over Wanda baseline on ARC-Challenge at 20% sparsity on Llama-3-8B-Instruct and Mistral-7B-Instruct. Evaluated on GSM8K, StrategyQA, and ARC-Challenge, CAP consistently outperforms alternatives at 10-20% sparsity. The key insight is that interventional (causal) measurement of head importance better captures functional contribution than correlational criteria — making pruned models that specifically preserve reasoning-critical attention structure.

■ *CAP offers a practical, training-free path to reducing production LLM inference costs by 10-20% while specifically protecting reasoning quality — the most business-critical capability. This is immediately applicable to any organization running Llama or Mistral variants in production.*

■ ■ THAILAND AI & TECH (1)

ADPT.news (via TechTalkThai RSS) | Thai | ■ MED | #Microsoft #AIAT #AI Engineering #Thailand

■ GLOBAL AI & TECH (2)

Beartai | Thai | ■ HIGH | #AI Safety #Japan #Banking #Cybersecurity

TechCrunch | International | ■ LOW | #Amazon #AWS #India #AI Infrastructure

Amazon has announced a fresh \$13 billion investment in AI infrastructure in India, significantly expanding its AWS cloud and data center presence in the country. This follows earlier India commitments and reflects a broader race by US tech giants to build AI capacity in South and Southeast Asia, where demand for cloud services and AI compute is growing rapidly. For Thailand and the ASEAN region, Amazon's India bet is directly relevant — large-scale infrastructure investment in India often precedes similar moves into Southeast Asia, as companies use the Indian buildout to refine regional operations. The investment will fund new AWS regions, AI-optimized data centers, and local AI talent development partnerships. India's AI cloud market is expected to become one of the world's largest within five years.

■ Amazon's \$13B India AI infrastructure bet signals that Southeast Asia is next — Thai businesses should expect dramatically improved AWS AI service availability and pricing within 24-36 months as this regional infrastructure comes online.