

[OpenAI] [GPT-5.6] [LLM] [Sol] [Terra] [Luna]

[OpenAI] [GPT-5.6] [■■■■■■■■■■] [AI Safety] [■■■■■]

[AI] [Future of Work] [Cory Doctorow] [████████] [██████████]

[Fable 5] [GPT-5.5] [Claude Opus] [benchmark] [coding] [FrontierCode]

Four Google DeepMind Senior Researchers Leave in Six Days for OpenAI and Anthropic

6 Google DeepMind 4 Noam Shazeer Transformer architecture co-lead Gemini OpenAI, John Jumper Nobel AlphaFold Anthropic Jonas Adler AI coding Google Alexander Pritzel pretraining AlphaFold Anthropic Bloomberg GPU run experiments...

Google DeepMind lost 4 top researchers in 6 days to OpenAI and Anthropic. The primary driver is compute allocation politics — a structural issue that risks compounding into ongoing talent exodus at Google.

[Google DeepMind] [Anthropic] [OpenAI] [talent] [Noam Shazeer] [John Jumper]

DeepSeek V4-Pro Price Cut Made Permanent at \$0.44/M Input Tokens — 5x Cheaper than GPT-5.5

DeepSeek V4-Pro 2026 0.44 input token 0.87 output token V4-Flash 0.14 0.28 GPT-5.5 (2.50/15) 5 input 17 output Gemini 3.5 Flash 3 sparse-attention hybrid 27% compute token 10% memory DeepSeek price floor Chinese open-weight model tier...

DeepSeek's permanent price floor at \$0.44/M input tokens creates a structural cost alternative for budget-sensitive AI applications. The 5x-17x price gap vs. Western frontier models is now large enough to drive meaningful enterprise workload migration.

[DeepSeek] [pricing] [inference cost] [GPT-5.5] [open-weight] [China AI]

GPT-5.6 'Kindle-Alpha' RC Nears Launch: 1.5M Context Window, 83% Polymarket Odds, Government Approval Customer-by-Customer

GPT-5.6 OpenAI 'kindle-alpha' AI release prediction market Polymarket 83% 28 2026 context window 1.5 token (1 GPT-5.5) UI front-end code Codex response reward audit pipeline 'Goblin Incident' GPT-5.5 OpenAI API...

GPT-5.6 is imminent but may face delayed API access due to government review. The automatic endpoint update (gpt-5.5-latest → GPT-5.6) is a critical risk for production systems not pinned to explicit versions.

[GPT-5.6] [OpenAI] [Codex] [release] [context window] [Polymarket]

OpenAI and Broadcom unveiled the Jalapeño custom AI inference chip designed to cut large-scale inference costs by approximately 50% compared to NVIDIA H200 GPUs. Built using HBM4 memory and fabricated at TSMC N2 node, the chip is purpose-built for transformer inference at OpenAI's scale. This marks OpenAI's first major move into custom silicon, following Google's TPU and Amazon's Trainium strategy. The partnership allows OpenAI to reduce dependency on NVIDIA while Broadcom gains a high-volume customer for its ASIC design services. At OpenAI's current inference volume, analysts estimate the chi...

[OpenAI] [Broadcom] [Jalapeño] [custom chip] [AI inference] [NVIDIA]

MAS-PromptBench Juyang Bai Laixi Shi benchmark
system prompt Multi-Agent LLM
Multi-Agent agent
system prompt
prompt agent multi-agent
prompt exponential workflow
task, communication protocol team size
optimize prompt ...

[multi-agent] [prompt-optimization] [LLM] [MAS] [benchmark] [arxiv]

Term	Frequency
UC Berkeley	10
Sergey Levine	5
sparse reward	25
dense process reward	10
Reinforcement Learning	15
RL	10
agent	15
discriminator	10
episode	5
reward	10
agent	15
optimal policy	10
manipulation	10
...	5

[reinforcement-learning] [reward-learning] [process-reward] [robotics] [Sergey-Levine] [Berkeley]

Memory Contagion: Cross-Temporal Propagation of Evaluator Bias via Agent Memory

AI agent
Memory Contagion bias
agent agent
bias
multi-agent
memory evaluation
AI agent LLM-as...

Memory Contagion evaluator bias agent memory
long-running AI agent systems

[agent-memory] [AI-safety] [LLM-evaluation] [evaluator-bias] [multi-agent] [arxiv]

Superhuman AI for Generals.io Using Self-Play Reinforcement Learning

AI superhuman Generals.io
real-time Self-Play Reinforcement Learning
AI Generals.io
imperfect information game
Chess Go
self-play RL AI
AI

Self-play RL superhuman performance Generals.io imperfect information real-time strategy game
proof-of-concept AI

[reinforcement-learning] [self-play] [game-AI] [superhuman] [generals-io] [incomplete-information]

AI Agents 101: The 2026 Stack | EP 607 | Daily AI News by GAI Insights

YouTube GAI Insights Episode 607 AI Agents 101: The 2026 Stack
AI agent 2026
agentic AI frameworks, orchestration layers, tool-use capabilities, memory systems multi-agent
coordination standard Qwen language
world models LLM
reasoning landscape
AI agent software ...

Episode 2026 AI Agent Stack Qwen language world models component
agentic AI deployment

[youtube] [AI-agents] [2026-stack] [GAI-insights] [agentic-AI] [Qwen]

Phasing Out Humans, AI Sin Eaters, Airplane Swarms and the AI Trolley Problem

Matt Wolfe ■■■■■■ FutureTools.io ■■■■■■
AI ■■■■■■
■■■■■■ AI Sin Eaters ■■■■■■ AI ■■■■■■
■■■■■■ airplane
swarms ■■■■■■ AI Trolley Problem
■■■■■■ AI
■■■■■■ AI
■■■■■■
■■■■■■

■ ■■■■■■ AI ■■■■■■ AI workforce displacement, Sin Eaters
framework, autonomous military systems ■■■■ trolley problem ■■■■■■ AI decision-making

[youtube] [matt-wolfe] [future-tools] [AI-jobs] [drone-swarms] [AI-ethics]

TECH TRENDS

1 ■■■■■■

Notion Mail ■■■■■■ 22 ■■■■■■ AI Agent
■■■■■■

Notion ■■■■■■ Notion Mail ■■■■■■
■■■■■■ 22 ■■■■■■ 2026
■■■■■■ Notion
■■■■■■ Notion Agent ■■■■■■ inbox
■■■■■■ AI Agent ■■■■■■ AI
Agent ■■■■■■
■■■■■■ Agentic AI
■■■■■■ productivity apps ■■■■■■

■ ■■■■ Notion Mail ■■■■■■ AI Agents ■■■■■■ productivity apps ■■■■■■
■■■■■■ Agentic AI paradigm

[Notion] [Notion Mail] [Notion Agent] [AI Agent] [Email]

THAILAND AI & TECH

1 ■■■■■■

■■■■■■ 3 Health Tech
■■■■■■

Techsauce ■■■■■■ 3 ■■■■■■ Health Tech ■■■■■■ AI
■■■■■■
■■■■■■ '■■■■■■'
■■■■■■ SPEAR-H
■■■■■■ AI ■■■■■■
■■■■■■
■■■■■■ AI healthcare ecosystem ■■■■■■
■■■■■■

■ ■■■■■■ Health Tech startups ■■■■■■ AI diagnostics ■■■■
■■■■■■ ecosystem ■■■■■■

[HealthTech] [■■■■] [SPEAR-H] [■■■■■■■■■■] [AI ■■■■■■]

Blognone · Thai

[OpenAI] [IPO] [■■■■■■■■■■] [Wall Street]

TechCrunch · International

[OpenAI] [Anthropic] [government] [AI policy] [GPT-5.6] [export control]

CNBC · International

[Anthropic] [Alibaba] [Qwen] [distillation attack] [AI security] [China]