

AI NEWS DAILY

2026-06-28

34 articles

AI Trends: 8

Tech: 3

Thailand: 15

Global: 8

12
HIGH URGENCY

13
MEDIUM

7
BREAKTHROUGHS

15
THAI SOURCES

10
INTERNATIONAL

5
RESEARCH

4
YOUTUBE

TOP INSIGHTS TODAY

It's not about Anthropic vs. OpenAI anymore

AI government approval frontier AI models

Apple Vision Pro exec is reportedly leaving for OpenAI

hardware Apple OpenAI wave AI competition physical hardware layer software API Apple

Asian AI startups launch Mythos-like models as Anthropic's export ban drags on

AI AI model AI AI AI

The fittest founder in the room got cancer. Here's how he used AI to fight back

LLM 'patient empowerment tool' use case

OpenAI limits GPT-5.6 rollout after government request, says restrictions shouldn't be the norm

Anthropic Claude 3.5 AI model 100% accurate in detecting AI-generated text. AI 100% accurate in detecting AI-generated text. AI 100% accurate in detecting AI-generated text.

Trump Admin releases Anthropic Mythos to be used by more than 100 US companies, agencies

Trump AI cybersecurity models 100% accurate in detecting AI-generated text. AI 100% accurate in detecting AI-generated text. AI 100% accurate in detecting AI-generated text.



AI TRENDS

8 ARTICLES

ARXIV CS.CL · RESEARCH

BREAKTHROUGH

HIGH

CAT-Q: Cost-efficient and Accurate Ternary Quantization for LLMs

Intel CAT-Q (Cost-efficient and Accurate Ternary Quantization) Post-training Quantization LLM Ternary 3-bit Training 512 Calibrate training tokens BitNet 100,000 Oral Paper ICML 2026 CAT-Q Learnable Modulation element-wise weight Scale-dependent L2 regularization weight ternary LLaMA-3 8B perplexity BitNet CAT-Q deploy LLM Low-bit LLM

arxiv research quantization LLM ICML2026

Actions

Developer:

- CAT-Q (BitTern) GitHub Intel Ternary
- CAT-Q LLM 7B-8B Accuracy Throughput

Business:

- CAT-Q Deploy LLM Edge Device IoT
- Ternary Quantization Intel ICML 2026 MLOps

ARXIV CS.CL · RESEARCH

BREAKTHROUGH

MED

Information-Aware KV Cache Compression for Long Reasoning

InfoKV KV Cache (Entropy-aware) Reasoning LLMs KV Cache Attention Score Reasoning chains InfoKV Forward Influence Metric Entropy Attention Pattern token Llama-3 DeepSeek-R1 baseline Multi-step Reasoning RAG memory budget InfoKV LLM Reasoning Memory document analysis

arxiv research KV-cache LLM reasoning

Actions

Developer:

- Forward Influence Long-context LLM Memory
- InfoKV DeepSeek-R1 Reasoning

Business:

- KV Cache Compression Memory Deploy Long-context LLM
- InfoKV RAG Document Processing

ARXIV CS.CL · RESEARCH BREAKTHROUGH HIGH

Nemotron-TwoTower: Diffusion Language Modeling with Pretrained Autoregressive Context

NVIDIA Research Nemotron-TwoTower Diffusion Language Model DLM Context Tower Autoregressive model context Diffusion Tower denoising pretrained Autoregressive model train Autoregressive 2.42 TwoTower Diffusion LM Production LLM Autoregressive Hybrid Diffusion-AR architecture

arxiv research diffusion LLM NVIDIA

Actions

Developer:

- Nemotron-TwoTower Hugging Face
- TwoTower Inference-efficient LLM

Business:

- Diffusion LLM Inference
- Throughput 2.42x

ARXIV CS.CL · RESEARCH

BREAKTHROUGH

HIGH

Empowering GUI Agents via Autonomous Experience Exploration and Hindsight Experience Utilization for Task Planning

PEEU (Planning Experience Exploration and Utilization) GUI Agent Multimodal MLLM Task Planning PEEU Autonomous Experience Exploration agent GUI Hindsight Experience Utilization training data TDHAF benchmark 7B PEEU Qwen2.5-VL-32B 4 ACL 2026 AI Agent Edge

arxiv research GUI-agent MLLM ACL2026

Actions

Developer:

- Hindsight Training Data Agent PEEU
- TDHAF Compositional Generalization Agent

Business:

- GUI Agent PEEU Automate
- ACL 2026 MLLM Agent Business Process Automation

ARXIV CS.LG · RESEARCH

MED

Tapered Language Models

Mila/Montreal Tapered Language Models (TLMs) LLM layer layer capacity (AR, MLM, Diffusion) learning layer layer Tapered LM capacity pattern Cosine Schedule MLP Width layer layer perplexity benchmark score parameter compute Taper MLP Width Cosine Schedule LLM

arxiv research LLM architecture efficiency

Actions

Developer:

- MLP Tapering Cosine Schedule Perplexi
- Paper Tapered LM Depth-aware Capacity Allocation

Business:

- Tapered LM Foundation Model
- Tapered Architecture Pre-train On-device

AI EXPLAINED · YOUTUBE

LOW

AI Explained — Latest AI Developments June 28, 2026

AI Explained AI 28 2026 GPT-5.6 White House Anthropic Mythos 5 ArXiv quantization diffusion LM AI development AI business leaders AI

youtube AI-news LLM analysis

Actions

Developer:

- AI Explained LLM Benchmark Capability Evaluation
- AI Safety Research

Business:

- AI Explained AI Policy
- Frontier Model AI Provider

WES ROTH · YOUTUBE

LOW

Wes Roth — AI News & Frontier Model Updates June 28, 2026

Wes Roth YouTuber AI frontier AI models 28 2026 OpenAI GPT-5.6 family Anthropic Mythos 5 US AI export control competitive landscape Roth two-tier AI market frontier

models [Google Gemini](#) [OpenAI GPT-4](#) Roth [Google Gemini](#) AI [Google Gemini](#) 6-12 [Google Gemini](#)

[youtube](#) [AI-news](#) [LLM](#) [frontier-AI](#)

▢ Actions

▢ **Developer:**

- [Google](#) Wes Roth [Google Gemini](#) AI Model Release [Google Gemini](#) Demo [Google Gemini](#)
- [Google](#) Benchmark Comparison [Google Gemini](#) LLM [Google Gemini](#)

▢ **Business:**

- [Google](#) Wes Roth [Google Gemini](#) AI [Google Gemini](#)
- [Google](#) AI Tool Announcements [Google Gemini](#) Workflow [Google Gemini](#)

YANNIC KILCHER · YOUTUBE

LOW

Yannic Kilcher — AI Research Paper Review June 2026

Yannic Kilcher AI Researcher [YouTube](#) [YouTube](#) [YouTube](#) Machine Learning [YouTube](#) Deep Learning [YouTube](#) 2026 [YouTube](#) [YouTube](#) annotate papers [YouTube](#) real-time [YouTube](#) [YouTube](#) papers [YouTube](#) ArXiv [YouTube](#) Tapered Language Models [YouTube](#) Mila [YouTube](#) [YouTube](#) depth-aware capacity allocation Nemotron-TwoTower [YouTube](#) NVIDIA [YouTube](#) diffusion LM [YouTube](#) CAT-Q [YouTube](#) Intel [YouTube](#) ternary quantization [YouTube](#) Kilcher [YouTube](#) [YouTube](#) [YouTube](#) ML engineers [YouTube](#) research frontier

[youtube](#) [AI-research](#) [paper-review](#) [deep-learning](#)

▢ Actions

▢ **Developer:**

- [Yannic Kilcher](#) [Yannic Kilcher](#) Paper [Yannic Kilcher](#) LLM [Yannic Kilcher](#) Diffusion Model [Yannic Kilcher](#)
- [Yannic Kilcher](#) Livestream Discussion [Yannic Kilcher](#)

▢ **Business:**

- [Yannic Kilcher](#) [Yannic Kilcher](#) AI [Yannic Kilcher](#)
- [Yannic Kilcher](#) Paper Review [Yannic Kilcher](#) AI Safety [Yannic Kilcher](#) Alignment [Yannic Kilcher](#) AI



TECH TRENDS

3 ARTICLES

TECHCRUNCH · INTERNATIONAL

⚡ BREAKTHROUGH

MED

The fittest founder in the room got cancer. Here's how he used AI to fight back

Connor Christou [founded](#) startup [founded](#) 34 [years](#) [ago](#) [when](#) [he](#) [was](#) [diagnosed](#) [with](#) [non-Hodgkin's lymphoma](#) [in](#) [2025](#) [his](#) [story](#) [with](#) [Claude](#) [from](#) [Anthropic](#) [as](#) [his](#) ["second opinion partner"](#) [helped](#) [him](#) [interpret](#) [his](#) [PET scan](#) [and](#) [blood markers](#) [using](#) [Claude](#) [with](#) [extended thinking](#) [mode](#) [to](#) [generate](#) [custom prompts](#) [for](#) [Claude](#) [to](#) [analyze](#) [his](#) [medical documents](#) [and](#) [reduce](#) [hallucination](#) [risk](#) [by](#) [prompting](#) [Claude](#) [to](#) [cite](#) [sources](#) [and](#) [provide](#) [confidence levels](#) [LLM](#) [as](#) [a](#) [patient empowerment tool](#) [for](#) [healthcare](#) [AI](#) [use](#) [case](#) [study](#)

[claude](#) [anthropic](#) [healthcare-ai](#) [medical-ai](#) [cancer](#)

Actions

Developer:

- [Claude](#) [API](#) [extended thinking mode](#)
- [Claude](#) [prompt](#) [LLM](#) [confidence level](#) [uncertainty](#)

Business:

- [AI](#) [health companion](#)
- [FDA/](#) [AI](#) [medical advisory tools](#)

TECHCRUNCH · INTERNATIONAL

LOW

Instagram is testing more ways to customize 'Your Algorithm'

Adam Mosseri [Instagram](#) [CEO](#) [announced](#) [that](#) [Instagram](#) [is](#) [testing](#) [new](#) [ways](#) [to](#) [customize](#) [its](#) ["Your Algorithm"](#) [in](#) [2025](#) [to](#) [improve](#) [user](#) [experience](#) [and](#) [engagement](#) [by](#) [introducing](#) [a](#) [new](#) ["Your Algorithm"](#) [section](#) [that](#) [will](#) [allow](#) [users](#) [to](#) [control](#) [what](#) [they](#) [see](#) [in](#) [their](#) [feed](#) [and](#) [swipe](#) [Reels](#) [Meta](#) [is](#) [testing](#) [a](#) [human-in-the-loop](#) [recommendation](#) [system](#) [to](#) [train](#) [AI](#) [models](#) [to](#) [better](#) [understand](#) [user](#) [intent](#) [and](#) [preferences](#) [social](#) [media](#) [platforms](#) [like](#) [TikTok](#) [and](#) [YouTube](#)

[instagram](#) [meta](#) [algorithm](#) [recommendation-system](#) [adam-mosseri](#)

Actions

Developer:

- ["explainable recommendation"](#) [user-controlled personalization](#)
- [feedback mechanism](#) [AI recommendation engine](#) [UX](#)

Business:

- [content strategy](#) [Instagram](#) [content](#) [explicitly](#)

- multi-factor authentication behavioral analytics

Business:

- Phishing AI
- cybersecurity AI

BEARTAI · THAI

HIGH

ATM/ AI

ATM

AI

Anthropic Claude Mythos 5

vulnerability

security

AI

business continuity planning

AI Safety Cybersecurity Japan Banking Anthropic

Actions

Developer:

- security AI
- anomaly detection AI

Business:

- business continuity AI
- cybersecurity solution AI

BEARTAI · THAI

MED

SITE 2026 : 'GDP

SITE 2026 (Startup x Innovation Thailand Expo 2026)

'GLOBAL INNOVATION IMPACT: THE YEAR OF INVESTMENT'

NIA (National Innovation Agency) SME Bank

keynote

GDP startups ecosystem

AI

tech hub

SITE 2026 NIA Innovation Thailand Startup

Actions

Developer:

- SITE 2026
-

Business:

-
- startup innovation NIA

BLOGNONE · THAI

LOW

AMD FSR Upscaling 4.1 Radeon RX 7000 (RDNA 3)

AMD Adrenalin Edition 26.6.2 FSR Upscaling 4.1 Radeon RX 7000 Series RDNA 3 FSR Upscaling 4.1 machine learning frame rate FSR 4.1 RDNA 4 (RX 9000 Series) RDNA 3 70 RX 7000 NVIDIA DLSS AMD PC gaming machine learning upscaling

AMD FSR GPU Radeon Machine Learning

Actions

Developer:

- FSR Upscaling 4.1 API AMD
- FSR 4.1 DLSS use case

Business:

- AMD GPU workstation AI-enhanced upscaling
- AI upscaling technology

BLOGNONE · THAI

HIGH

Anthropic Mythos 5

Anthropic Claude Mythos 5 Fable 5 cybersecurity guardrails bypass frontier AI models

Actions

Developer:

- AI
- compliance AI

Business:

- AI AI
- AI

[] Google Gemini Enterprise Meta

The Information 2026 Google Meta Gemini Enterprise Meta Google infrastructure Meta enterprise Gemini AI Google capacity cloud AI providers AI enterprise Google Cloud Microsoft Azure enterprise AI

Actions

Developer:

- AI
- SLA capacity Cloud

Business:

- AI roadmap infrastructure
- multi-cloud strategy AI

Notion Mail 22 AI Agent

Notion Notion Mail 22 2026 AI Agent email client AI Agent productivity tools traditional human interface Notion Notion AI email agents AI Agent disruption SaaS productivity

Actions

Developer:

- AI Agent capabilities
- AI Agent adoption product roadmap

Business:

- workflow AI Agent
- AI Agent platform

OpenAI GPT-5.6 Sol

OpenAI GPT-5.6 AI Sol reasoning mode 'max' 'ultra' subagents 5 input tokens, Terra Sol Luna realtime OpenAI White House AI

Actions

Developer:

- Sol, Terra Luna use case
- API GPT-5.6 migrate

Business:

- GPT-5.6
- use case Sol

Paul Meade Vision Pro Apple

Mark Gurman Bloomberg Paul Meade Apple Vision Pro headset smart glasses OpenAI AI hardware device Jony Ive Chief Design Officer Apple Apple strategy Vision Pro smart glasses Meta Ray-Ban Meade OpenAI hardware war for talent AI hardware

Actions

Developer:

- OpenAI Hardware ecosystem
- AI-first hardware design

Business:

- AI OpenAI
- AI Hardware

Qualcomm Dragonfly Data Center AI

Qualcomm Dragonfly Data Center AI Dragonfly C1000 Qualcomm Oryon Snapdragon X Elite AI inference workloads Qualcomm edge data center telco cloud power efficiency performance AI inference Dragonfly C1000 x86 processors Qualcomm Data Center AI ARM-based processors Intel AMD

Actions

Developer:

- SDK Qualcomm Dragonfly AI300
- performance benchmark Dragonfly C1000 2028

Business:

- Qualcomm Data Center AI
- roadmap pricing Dragonfly NVIDIA AMD

OpenAI GPT-5.6

CNN White House OpenAI GPT-5.6 pre-announcement restriction model Anthropic Mythos 5

OpenAI framework frontier AI models AI

OpenAI Whitehouse AI Policy GPT-5.6 USA

Actions

Developer:

- API AI
- AI

Business:

- AI
- AI

TECHSAUCE · THAI

MED

12,000 AI,

12,000 2025-2026 AI, chip design 10,000 Tsinghua Peking University AI Engineering Made in China 2035 AI talent pool

China AI Education Robotics Policy

Actions

Developer:

- AI
- AI

Business:

- AI talent acquisition development AI
- upskilling AI

TECHSAUCE · THAI

MED

1 SDG 3 'Holistic Wellbeing'

1 SDG 3 (Good Health and Well-being) 93.6 THE Sustainability Rankings 2026 'Holistic Wellbeing' medical AI, genomics, precision medicine digital health wellness economic opportunity healthcare cost Wellness Economy hub health technology medical

Mahidol SDG3 Medical AI Wellness Economy Thailand

Actions

Developer:

-  Medical AI solution 
- open datasets AI

Business:

-  Wellness Economy  Medical AI  HealthTech
- Wellness Economy Medical AI HealthTech

TECHSAUCE · THAI

LOW

Data Center **NVIDIA** **45**

NVIDIA AI Factory GPU 'Rubin' liquid cooling 45°C heating PUE (Power Usage Effectiveness) data center AI workloads sustainable AI infrastructure

NVIDIA AI Data Center Cooling GPU

□ Actions

Developer:

- ☐ liquid cooling technology ☐ AI workload ☐ infrastructure
- ☐ warm water cooling ☐ air cooling ☐

Business:

- ☐ sustainable cooling strategy ☐ AI Data Center ☐
- ☐ warm water cooling ☐

TECHSAUCE · THAI

MED

AI-first Operating System

World Economic Forum Summer Davos 2026 Techsauce AI-first Operating System AI operations tool embedment AI process customer service supply chain WEF AI realtime decision-making

AI-first OS enterprise AI

AI WEF Summer Davos AI-first Enterprise

Actions

Developer:

- AI-first organization real-time data pipeline AI decision s
- AI integration framework AI-first

Business:

- AI
- AI upskilling AI-first

GLOBAL AI & TECH 8 ARTICLES

TECHCRUNCH · INTERNATIONAL

HIGH

It's not about Anthropic vs. OpenAI anymore

Russell Brandom AI TechCrunch AI regulation Anthropic OpenAI AI model Trump AI cybersecurity biorisk AI model actor access collective action independent review bodies

openai anthropic trump ai-regulation frontier-ai

Actions

Developer:

- AI NIST, CISA Commerce Dept. review AI models
- AI compliance requirements 2026-2027

Business:

- AI supply chain diversification frontier model
- independent AI review bodies standard-setting

TECHCRUNCH · INTERNATIONAL

BREAKTHROUGH

MED

Apple Vision Pro exec is reportedly leaving for OpenAI

Paul Meade Apple Vision Pro headset smart glasses AI Apple

hardware OpenAI Bloomberg Apple strategy smart glasses Meta Ray-Ban Vision Pro wave AI competition physical hardware layer OpenAI Jony Ive CDO Apple AI device interact AI Apple

apple openai vision-pro hardware paul-meade

Actions

Developer:

- OpenAI/Jony Ive AI device SDK developer ecosystem
- Apple smart glasses API apps platform

Business:

- AI wearable/hardware strategy Meta Ray-Ban Apple/Ope
- launch AI device OpenAI paradigm interact AI

Asian AI startups launch Mythos-like models as Anthropic's export ban drags on

AI Trump Anthropic Mythos 5 Fable 5 Sakana AI Google Brain Fugu multi-agent orchestration model Mythos 5 cybersecurity 360 AI CyberShield Pro Mythos 5 code AI model AI

asia anthropic export-ban mythos sakana-ai

Actions

Developer:

- Sakana Fugu API orchestration
- multi-vendor AI export controls

Business:

- AI supply chain US AI providers
- AI

Indian payments chief thinks AI will be heavily involved in next era of digital payment growth

Dilip Asbe, CEO of NPCI (National Payments Corporation of India), says UPI will be the primary mode of payment for 750 million Indians. He says 24-hour AI-powered small language models (SLMs) will be used for UPI transactions, on-device SLMs will be used for voice interface, and AI will be used for cross-border payments, compliance, and fraud detection.

india npci upi digital-payments fintech

Actions

Developer:

- on-device SLM frameworks like MediaPipe LLM Inference or llama.cpp for fintech applications
- voice payment interface

Business:

- UPI transactions, PromptPay, e-Wallet
- cross-border payment, UPI, AI middleware

TECHCRUNCH · INTERNATIONAL

MED

OpenAI poaches Uber India chief to lead its biggest market outside the US

OpenAI has poached Prabhjeet Singh, CEO of Uber India and South Asia, to lead its biggest market outside the US. Singh will be managing the company's operations in India, Asia Pacific, and other regions. He will be responsible for consumer growth, enterprise adoption, partnerships, regulatory engagement, and operations. OpenAI has a long history of partnerships with companies like Reliance JioHotStar, Tata Group, and Anthropic. It also has a presence in AI labs and the Global South.

openai india uber prabhjeet-singh asia-pacific

Actions

Developer:

- OpenAI AI deployment engineer, solutions engineer
- OpenAI partnership ecosystem, Asia Pacific, API integration

Business:

- OpenAI India playbook partnership opportunities
- OpenAI startup program credits support Asia Pacific expansion

TECHCRUNCH · INTERNATIONAL

HIGH

OpenAI limits GPT-5.6 rollout after government request, says restrictions shouldn't be the norm

OpenAI announced GPT-5.6 rollout after government request, saying restrictions shouldn't be the norm. The rollout is being limited to Sol (), Terra (), and Luna () models. OpenAI is also limiting the 'reasoning mode' to 'max' and 'ultra' modes, and subagents. The rollout is being limited to 5 input tokens and 30 output tokens. Terra is also being limited to OpenAI's cost-performance. OpenAI is also limiting the rollout to AI models.

openai gpt-5.6 trump government ai-regulation

Actions

Developer:

- trusted partner OpenAI GPT-5.6 Sol
- GPT-5.6 Terra Luna cost-performance

Business:

- AI frontier AI models
- AI strategy frontier models

TECHCRUNCH · INTERNATIONAL

MED

SoftBank's CEO isn't the only one with questions about Elon Musk's orbital data center hype

Masayoshi Son SoftBank Elon Musk data center Musk early stage TechCrunch Musk SpaceX orbital data center launch capacity SpaceX AI infrastructure compute

softbank elon-musk spacex orbital-data-center ai-infrastructure

Actions

Developer:

- cloud compute Groq Baseten inference
- compute providers

Business:

- compute provider
- custom AI chip Jalapeño OpenAI inference

Trump Admin releases Anthropic Mythos to be used by more than 100 US companies, agencies

Anthropic Mythos 5 Fable 5 cybersecurity Trump Anthropic Mythos 5 100 bypass guardrails cybersecurity AI cybersecurity models AI

trump anthropic mythos-5 fable-5 export-ban

Actions

Developer:

- Mythos 5
- fallback strategy open-source cybersecurity models Meta Llama 3 Cyber Anthropic models

Business:

- frontier AI
- Asian AI alternatives Baidu, Alibaba Qwen open-source LLMs cybers cases