



⚡ AI NEWS DAILY

📅 2026-06-29 | Generated: 2026-06-29T03:55 ICT | Pipeline v2.0



TOP INSIGHTS – HIGH URGENCY

HIGH

Introducing computer use in Gemini 3.5 Flash

Computer use in Gemini 3.5 Flash enables agentic automation for tasks like web browsing, data analysis, and code execution. Anthropic's browser automation capabilities are a key differentiator.

Agentic AI

Computer Use

RPA

HIGH

Previewing GPT-5.6 Sol: a next-generation model series

OpenAI's GPT-5.6 Sol series introduces a new architecture with improved cache efficiency and performance. It targets both developer and enterprise markets.

LLM

API

Cost Optimization

HIGH

OpenAI previewing GPT-5.6 Sol: prompt injection resistance and new pricing

GPT-5.6 Sol features enhanced prompt injection resistance and competitive pricing. It also includes agentic AI capabilities for complex tasks.

AI Safety

Prompt Injection

Security

HIGH

When the Trump administration cracks down on Anthropic, who benefits?

Export control on Anthropic could impact its competitive advantage in the AI market. This raises geopolitical risks for the AI industry.

AI Policy

Export Control

Geopolitics

HIGH

CAT-Q: Cost-efficient and Accurate Ternary Quantization for LLMs

Ternary quantization (CAT-Q) improves LLM performance on edge devices and consumer hardware. It reduces inference cost while maintaining accuracy.

LLM Quantization Model Compression

HIGH

Empowering GUI Agents via Autonomous Experience Exploration and Hindsight Experience Utilization for Task Planning

hindsight learning GUI agent human annotation automation workflow

GUI Agents Agents LLM

AI TRENDS

18 articles

Google DeepMind

MEDIUM

< BREAKTHROUGH

DiffusionGemma: 4x faster text generation

DiffusionGemma: 4x faster text generation

Google DeepMind DiffusionGemma diffusion autoregressive approach LLM token latency use case real-time response AI transformer --- DiffusionGemma autoregressive LLM inference cost latency DiffusionGemma latency-sensitive applications real-time chat output quality diffusion approach autoregressive model diffusion text models inference cost use case 4x speed improvement user experience medium — LLM Architecture, Diffusion Models, Inference Speed

LLM Architecture Diffusion Models Inference Speed Research

Tech Impact: DiffusionGemma autoregressive LLM in

Dev Actions:

- DiffusionGemma latency-sensitive applications real-time chat
- output quality diffusion approach autoregressive model

Business Actions:

- diffusion text models inference cost
- use case 4x speed improvement user experience

Source

2026-06-01

Google DeepMind

HIGH

➤ BREAKTHROUGH

Introducing computer use in Gemini 3.5 Flash

Google DeepMind 发布了 Computer Use 的 Gemini 3.5 Flash

Google DeepMind 发布了 computer use 的 Gemini 3.5 Flash，
这是 AI 领域的一个重大突破。Anthropic 的 Claude 模型
Gemini 3.5 Flash 在 agentic AI 方面表现出色，
能够处理 real-world 任务。Flash 模型
使用 Google 的 computer use 能力，
开发者 --- Computer use 的 Gemini 3.5 Flash
能够进行 agentic automation。Anthropic 的
browser automation 的 Gemini 3.5 Flash
computer use API 的 RPA 的 browser automation use cases
的 accuracy 的 cost 的 Gemini 3.5 Flash 的 Claude computer use
use cases 的 computer use AI 的 manual desktop workflow 的
security policy 的 AI agent 的

Agentic AI Computer Use RPA Browser Automation

Tech

Impact: Computer use 的 Gemini 3.5 Flash 的 agentic automation 的 browser automation

Dev Actions:

- Gemini 3.5 Flash computer use API 的 RPA 的 browser automation use cases
- accuracy 的 cost 的 Gemini 3.5 Flash 的 Claude computer use

Business Actions:

- use cases 的 computer use AI 的 manual desktop workflow 的
- security policy 的 AI agent 的

Source

2026-06-01

Previewing GPT-5.6 Sol: a next-generation model series

OpenAI 发布了 GPT-5.6 Sol 系列模型

OpenAI 发布了 GPT-5.6 系列模型，包括 limited preview 模型 Sol、Terra 和 Luna。Sol 是 flagship model，价格为 \$5 input / \$30 output，支持 1M tokens。Terra 价格为 \$2.50 / \$15，Luna 价格为 \$1 / \$6。GPT-5.6 系列模型支持 prompt caching，具有 explicit cache breakpoints，minimum cache life 30 秒，cache writes 1.25x input，cache reads 90%。OpenAI 表示，GPT-5.6 系列模型将用于 trusted partners，包括 OpenAI 的 GPT-5.6 系列模型 Sol/Terra/Luna，cache 支持 developer 和 enterprise 使用。GPT-5.6 Terra 模型支持 GPT-5.5 的 2 倍 cache breakpoints，production 使用 token 使用 tier 使用 case 使用

LLM API Cost Optimization Caching

Tech Impact: OpenAI 发布了 GPT-5.6 系列模型 Sol/Terra/Luna，cache 支持

Dev Actions:

- GPT-5.6 Terra 模型支持 GPT-5.5 的 2 倍 cache breakpoints
- cache breakpoints production token 使用

Business Actions:

- tier 使用 case 使用
- trusted partners limited preview 使用

Simon Willison

MEDIUM

GLM-5.2 is probably the most powerful text-only open weights LLM

GLM-5.2 is probably the most powerful open weights LLM for text-only tasks. Simon Willison, co-founder of Z.ai, has announced the release of open weights for GLM-5.2 under a MIT license. The model has 753 billion parameters, uses a Mixture of Experts (MoE) architecture with 40 active parameters, and has a context window of 1 million tokens. This is a significant improvement over GLM-5.1, which had a context window of 200,000 tokens. Willison notes that GLM-5.2 is an open weights LLM that can handle visual input, unlike Z.ai's previous release which was under a MIT license and a closed model. He also mentions that GLM-5.2 is an open weights model under a MIT license, which is a significant improvement over frontier models like GLM-5.2 and OpenRouter. He also mentions that GLM-5.2 has performance that is on par with frontier models, and that it can be self-hosted on GLM-5.2 MoE infrastructure. Finally, he mentions that GLM-5.2 is a significant improvement over US-based AI providers.

Open Source AI LLM MoE Architecture China AI

Tech Impact: GLM-5.2 is an open weights model under a MIT license, which is a significant improvement over frontier models.

Dev Actions:

- Evaluate GLM-5.2's performance on OpenRouter compared to frontier models.
- Explore self-hosting GLM-5.2 MoE infrastructure.

Business Actions:

- Assess GLM-5.2's performance against US-based AI providers.
- Consider compliance and data residency requirements for model deployment.

[Source](#)

2026-06-17

Simon Willison

HIGH

OpenAI previewing GPT-5.6 Sol: prompt injection resistance and new pricing

Simon Willison [@simonwillison](#) GPT-5.6 Sol: [@simonwillison](#) prompt injection [@simonwillison](#)

Simon Willison [@simonwillison](#) AI [@simonwillison](#) developer tools [@simonwillison](#) [@simonwillison](#) [@simonwillison](#) GPT-5.6 Sol [@simonwillison](#) prompt injection resistance [@simonwillison](#) system card [@simonwillison](#) OpenAI [@simonwillison](#) Willison [@simonwillison](#) Fernando Irarrázaval [@simonwillison](#) challenge [@simonwillison](#) 2,000 [@simonwillison](#) AI assistant [@simonwillison](#) email [@simonwillison](#) Opus 4.6 [@simonwillison](#) 6,000 [@simonwillison](#) [@simonwillison](#) AI labs [@simonwillison](#) training [@simonwillison](#) frontier models [@simonwillison](#) injection attack [@simonwillison](#) --- GPT-5.6 Sol [@simonwillison](#) prompt injection resistance [@simonwillison](#) [@simonwillison](#) agentic AI [@simonwillison](#) GPT-5.6 Sol [@simonwillison](#) prompt injection scenarios [@simonwillison](#) input [@simonwillison](#) external sources [@simonwillison](#) GPT-5.6 system card [@simonwillison](#) improved injection resistance [@simonwillison](#) deploy agentic system [@simonwillison](#) security policy [@simonwillison](#) AI deployment [@simonwillison](#) best practices [@simonwillison](#) system card [@simonwillison](#)

AI Safety

Prompt Injection

Security

Agentic AI

Tech Impact: GPT-5.6 Sol [@simonwillison](#) prompt injection resistance [@simonwillison](#)

Dev Actions:

- [@simonwillison](#) GPT-5.6 Sol [@simonwillison](#) prompt injection scenarios [@simonwillison](#) input [@simonwillison](#) external sources
- [@simonwillison](#) GPT-5.6 system card [@simonwillison](#)

Business Actions:

- [@simonwillison](#) improved injection resistance [@simonwillison](#) deploy agentic system [@simonwillison](#)
- [@simonwillison](#) security policy [@simonwillison](#) AI deployment [@simonwillison](#) best practices [@simonwillison](#) system card [@simonwillison](#)

[Source](#)

2026-06-26

GPT-5.2 first impressions: a powerful update, especially for business tasks and workflows

📄 GPT-5.2: 📄

VentureBeat 📄 early testers 📄 GPT-5.2 📄 OpenAI 📄 public release 📄 deep autonomous reasoning 📄 coding 📄 workflow 📄 incremental 📄 dramatic 📄 pattern 📄 LLM 📄 use case 📄 power users 📄 enterprise 📄 AI 📄 enterprise-focused --- GPT-5.2 📄 trend 📄 LLM 📄 enterprise reasoning 📄 general chat 📄 GPT-5.2 📄 complex autonomous reasoning tasks 📄 improvement 📄 GPT-5.2 📄 Claude Tag 📄 Gemini 3.5 📄 multi-step workflow 📄 ROI 📄 GPT-5.2 📄 use case 📄 business automation 📄 proof of concept 📄 deep reasoning 📄 migrate 📄: medium — 📄 LLM, Enterprise AI, Reasoning

LLM Enterprise AI Reasoning Business Automation

Tech Impact: GPT-5.2 📄 trend 📄 LLM 📄 enterprise reasoning 📄 general

Dev Actions:

- 📄 GPT-5.2 📄 complex autonomous reasoning tasks 📄 improvement 📄
- 📄 GPT-5.2 📄 Claude Tag 📄 Gemini 3.5 📄 multi-step workflow

Business Actions:

- 📄 ROI 📄 GPT-5.2 📄 use case 📄 business automation
- 📄 proof of concept 📄 deep reasoning 📄 migrate

CAT-Q: Cost-efficient and Accurate Ternary Quantization for LLMs

CAT-Q: Ternary Quantization Ternary LLMs (ICML 2026 Oral)

CAT-Q model compression ICML 2026 machine learning quantization ternary (3-level: -1, 0, +1) LLM CAT-Q ternary quantization cost function compression ratio accuracy preservation checkpoints GitHub benchmark CAT-Q memory requirement full-precision model --- Ternary quantization ICML Oral deploy LLM edge device consumer hardware inference cost CAT-Q GitHub (IntelChina-AI/BitTern) CAT-Q GPTQ AWQ throughput perplexity hardware cloud inference cost ternary-quantized model full-precision

LLM Quantization Model Compression Efficient AI

Tech Impact: Ternary quantization ICML Oral deploy LLM edge device consumer hardware inference cost

Dev Actions:

- CAT-Q GitHub (IntelChina-AI/BitTern)
- CAT-Q GPTQ AWQ throughput perplexity hardware

Business Actions:

- cloud inference cost ternary-quantized model full-precision
- on-premise deployment LLM quantization GPU

[Source](#)

2026-06-26

Empowering GUI Agents via Autonomous Experience Exploration and Hindsight Experience Utilization for Task Planning

Empowering GUI Agents via Autonomous Experience Exploration and Hindsight Experience Utilization (ACL 2026)

GUI agent Empowering AI via Autonomous Experience Exploration and Hindsight Experience Utilization for Task Planning agentive AI ACL 2026 Autonomous Experience Exploration Hindsight Experience Utilization agent hindsight agent feedback agent training data human annotation benchmark GUI planning baseline computer-use agent --- hindsight learning GUI agent human annotation automation workflow hindsight experience replay computer-use agent RPA autonomous exploration pipeline training data GUI task GUI agent manual process back-office computer-use AI automation 12 high — GUI Agents, Agents, LLM

GUI Agents

Agents

LLM

Automation

Tech Impact: hindsight learning GUI agent human annotation workflow

Dev Actions:

- hindsight experience replay computer-use agent RPA
- autonomous exploration pipeline training data GUI task

Business Actions:

- GUI agent manual process back-office
- computer-use AI automation 12

Information-Aware KV Cache Compression for Long Reasoning

Information-Aware KV Cache Compression for Long Reasoning

reasoning long-chain LLM KV cache memory Information-Aware KV Cache Compression token information-theoretic key-value pair attention-score-based compression token attention weight long-form reasoning benchmark KV cache size reasoning reasoning model o3 DeepSeek-R2 chain-of-thought --- KV cache reasoning model memory budget deploy LLM reasoning production information-aware KV cache compression reasoning workload chain-of-thought KV cache compression inference engine vLLM SGLang GPU memory KV cache compression reasoning-heavy application

LLM Reasoning KV Cache Inference Optimization

Tech Impact: KV cache reasoning model reasoning production

Dev Actions:

- information-aware KV cache compression reasoning workload chain-of-thought
- KV cache compression inference engine vLLM SGLang

Business Actions:

- GPU memory KV cache compression reasoning-heavy application
- scaling AI reasoning memory efficiency

MAGE-RAG: Multigranular Adaptive Graph Evidence for Agentic Multimodal RAG in Long-Document QA

MAGE-RAG: Agentic Multimodal RAG Long-Document QA

MAGE-RAG Retrieval-Augmented Generation Agentic (multimodal) Multigranular Adaptive Graph Evidence agent Long-Document Question Answering benchmark baseline RAG enterprise document processing research assistant --- MAGE-RAG RAG document intelligence multigranular graph construction RAG pipeline integration MAGE-RAG long-context LLM document QA multimodal RAG knowledge management system graph-based evidence: medium — RAG, Multimodal AI, Agents

RAG Multimodal AI Agents LLM

Tech Impact: MAGE-RAG RAG

Dev Actions:

- multigranular graph construction RAG pipeline
- integration MAGE-RAG long-context LLM document QA

Business Actions:

- multimodal RAG knowledge management system
- graph-based evidence

MAS-PromptBench: When Does Prompt Optimization Improve Multi-Agent LLM Systems?

MAS-PromptBench: Prompt Optimization Multi-Agent LLM Systems

multi-agent LLM system paradigm AI application prompt optimization MAS-PromptBench framework prompt multi-agent agent prompt optimization task agent topology communication agent configuration prompt optimization optimize multi-agent pipeline --- prompt optimization multi-agent system MAS-PromptBench baseline evaluate prompt optimization strategy multi-agent pipeline agent topology prompt optimization prompt engineering multi-agent benchmark budget prompt optimization Multi-Agent Systems, LLM, Prompt Engineering

Multi-Agent Systems

LLM

Prompt Engineering

Agents

Tech Impact: prompt optimization

Dev Actions:

- MAS-PromptBench baseline evaluate prompt optimization strategy multi-agent pipeline
- agent topology prompt optimization

Business Actions:

- prompt engineering multi-agent benchmark
- budget prompt optimization

[Source](#)

2026-06-23

On-Policy Distillation with Curriculum Turn-level Guidance for Multi-turn Agents

Distillation On-Policy Curriculum Turn-level Guidance Multi-turn Agents

LLM agent multi-turn agent turn On-Policy Distillation Curriculum Turn-level Guidance agent teacher model turn trajectory Curriculum turn agent benchmark agentic task AI agent customer service coding assistant --- agent training multi-turn agent reward model curriculum-based distillation coding agent customer service agent turn-level guidance trajectory-level reward RLHF multi-turn agent RLHF agent on-policy distillation chatbot: medium — LLM, Agents, Reinforcement Learning

LLM Agents Reinforcement Learning Distillation

Tech Impact: agent training multi-turn agent

Dev Actions:

- curriculum-based distillation coding agent customer service agent
- turn-level guidance trajectory-level reward RLHF

Business Actions:

- multi-turn agent RLHF
- agent on-policy distillation chatbot

[Source](#)

2026-06-16

OPID: On-Policy Skill Distillation for Agentic Reinforcement Learning

OPID: On-Policy Skill Distillation 提高 Agentic Reinforcement Learning

提高 agent 训练 reinforcement learning 效率 sample efficiency
稀疏 reward signal 稀疏 sparse 提高 agent 训练 OPID (On-Policy Skill Distillation) 提高 skill knowledge 专家
teacher policy 提高 agent 训练 OPID
distillation 提高 on-policy skill 分布
distribution 提高 state 提高 agent 训练 off-policy skill
静态 static dataset 分布 shift 提高
imitation learning 提高 agent 训练 agentic task
OPID 提高 final performance 提高 baseline
RL 提高 --- OPID 提高 agent 训练 sample
compute 提高 agentic system 提高 on-policy skill distillation
bootstrap 提高 RL agent 提高 domain
提高 OPID 提高 RLHF 提高 GRPO pipeline 提高 convergence 提高 LLM agent

Reinforcement Learning Agents LLM Skill Learning

Tech Impact: OPID 提高 agent 训练 sample 提高 提高 c

Dev Actions:

- 提高 on-policy skill distillation 提高 bootstrap 提高 RL agent 提高 domain
- 提高 OPID 提高 RLHF 提高 GRPO pipeline 提高 convergence 提高 LLM agent

Business Actions:

- 提高 expert demonstration data 提高 domain 提高 skill distillation
- 提高 ROI 提高 skill distillation 提高 agent 提高 pure RL 提高 compute

[Source](#)

2026-06-26

Tapered Language Models

Tapered Language Models: Tapered Architecture

Tapered Language Models LLM layer input layer output layer taper transformer inference memory footprint computation layer tapered architecture parameter efficient LLM --- Tapered architecture LLM compute memory Tapered LM task non-uniform layer size optimize inference latency tapered model cloud inference cost upgrade LLM: low — LLM, Architecture, Efficient AI

LLM Architecture Efficient AI Inference Optimization

Tech Impact: Tapered architecture LLM compute mem

Dev Actions:

- Tapered LM task
- non-uniform layer size optimize inference latency

Business Actions:

- tapered model cloud inference cost
- upgrade LLM

DeepSeek Releases DSpark, a Speculative Decoding Framework That Accelerates DeepSeek-V4 Per-User Generation 60-85% Over MTP-1

DeepSeek 发布了 DSpark: Framework Speculative Decoding 框架，可将推理速度提升 85%。

DeepSeek 发布了 DSpark 框架，该框架基于 speculative decoding 技术，旨在加速 DeepSeek-V4 模型的推理速度。DSpark 框架通过引入 draft module 和 target model，利用 token 的冗余性进行 speculative decoding。DSpark 框架采用 parallel draft backbone 和 sequential Markov head，结合 suffix decay 和 parallel drafter，实现 confidence-scheduled verification。该框架在 token 生成过程中，通过 GPU 加速，可将生产速度提升 60-85%。相比 V4-Flash，DSpark 可将推理延迟降低 57-78%。DeepSeek-V4-Pro 使用 DSpark 框架，可将 LLM 推理延迟降低 85%，同时保持模型权重不变，提升用户体验，降低 API 成本。DeepSpec 框架已在 GitHub 开源，支持 draft model 和 target model 的集成。DSpark-5 checkpoint 可在 DeepSeek-V4 模型上使用，实现 latency improvement。该框架在生产环境中，可显著提升 ROI，降低 GPU 成本。

LLM Inference Optimization Speculative Decoding DeepSeek

Tech Impact: DSpark 可将推理延迟降低 85%，提升模型权重，提升用户体验，降低 API 成本。

Dev Actions:

- 集成 DeepSpec 框架，支持 draft model 和 target model 的集成。
- 使用 DSpark-5 checkpoint 可在 DeepSeek-V4 模型上使用，实现 latency improvement。

Business Actions:

- 提升 ROI，降低 DSpark-based serving 的 GPU 成本。
- 集成 open-source DeepSpec 框架，支持 draft model 的集成。

[Source](#)

2026-06-27

2026-06-26

🔗 Claude Tag 🤖 AI 📧 Slack 📌

📌 Claude Tag 🤖 AI 📧 Slack 📌 @Claude

Anthropic 📌 'Claude Tag' 📌 Slack 📌 @Claude 📌 Automated 📌 AI 📌 'AI Teammate' 📌 Claude Tag 📌 Agentic 📌 AI 📌 Multi-step 📌 AI 📌 Workflow --- Claude Tag 📌 Slack 📌 Messaging App 📌 AI-powered Work Hub 📌 LLM 📌 Claude API 📌 Slack Integration 📌 Custom Workflow Automation 📌 Agentic AI System 📌 Multi-step Task 📌 Chat Platform 📌 Claude Tag 📌 Customer Support 📌 Project Management 📌 Upskill 📌 AI Agents: high — LLM, Agentic AI, Enterprise Software

LLM Agentic AI Enterprise Software Slack

Tech Impact: Claude Tag 📌 Slack 📌 Messaging App 📌 AI-powered Work Hub 📌

Dev Actions:

- 📌 Claude API 📌 Slack Integration 📌 Custom Workflow Automation 📌
- 📌 Agentic AI System 📌 Multi-step Task 📌 Chat Platform

Business Actions:

- 📌 Claude Tag 📌 Customer Support 📌 Project Management 📌
- 📌 Upskill 📌 AI Agents 📌

AI-first Operating System

AI Strategy	Digital Transformation	Enterprise AI	WEF
-------------	------------------------	---------------	-----

Tech Impact: ☒ Redesign ☐ AI-first ☐

- Data Pipeline AI Infrastructure Real-time AI Models
- AI Orchestration Layer AI Models AI Models

- ☐ AI Readiness ☐ Transformation Roadmap ☐
12-18 ☐
- ☐ Chief AI Officer ☐ AI Center of Excellence ☐
☐ AI-first Organization

2026-06-23

Introducing North Mini Code: Cohere's first model for developers

Cohere introduces North Mini Code, a new developer-focused model.

Cohere introduces North Mini Code, a new coding model designed to assist developers in practical software engineering tasks. It achieves a high benchmark score in agentic coding tasks, such as generating code from requirements, debugging, and code review. Cohere's North Mini Code is an enterprise NLP model that integrates with existing coding tools market. It is the first model from Cohere's HuggingFace community blog, demonstrating a strong community presence. Cohere is open source, and North Mini Code is also open source. Cohere's enterprise NLP model is designed to assist developers in enterprise-grade coding tasks. North Mini Code is a real coding task model that can be used alongside Claude Code and GitHub Copilot. API pricing, enterprise license terms, and North Mini Code team usage are also available.

[Coding AI](#) [Developer Tools](#) [Enterprise AI](#) [Agentic Coding](#)

Tech Impact: North Mini Code is Cohere's first enterprise NLP model for coding tools, designed to assist developers in enterprise-grade coding tasks.

Dev Actions:

- North Mini Code is a real coding task model that can be used alongside Claude Code and GitHub Copilot.
- API pricing, enterprise license terms, and North Mini Code team usage are also available.

Business Actions:

- North Mini Code is a developer tools stack that integrates with Cohere's existing tools.
- Security and compliance features are also available in GitHub Copilot Enterprise.

[Source](#)

2026-06-09

HuggingFace

MEDIUM

Introducing Serge: GitHub-Native AI Code Review

HuggingFace introduces Serge, an AI Code Review tool natively integrated with GitHub.

HuggingFace introduces Serge, an AI code review tool natively integrated with GitHub. This native integration allows developers to configure the service directly within their GitHub ecosystem, simplifying the setup process. Serge is designed to streamline the code review workflow, providing a seamless experience for developers. As an external service, Serge integrates with the existing developer workflow, leveraging AI to enhance the review process. This tool represents a significant trend in AI-powered development tools, moving from standalone products to native integrations within popular workflows. Serge natively integrates with the GitHub workflow, providing an external service that enhances security for developer teams. By integrating with the GitHub repository, Serge provides quality review feedback, similar to GitHub Copilot code review and CodeRabbit. Serge's code review capabilities are designed to be secure, adhering to data privacy policies and maintaining a secure codebase.

Developer Tools

Code Review

GitHub

AI Coding

Tech Impact: Serge introduces native AI code review into the GitHub workflow, providing an external service that enhances security for developer teams.

Dev Actions:

- Integrate Serge with the GitHub repository to enhance quality review feedback.
- Utilize Serge for GitHub Copilot code review and CodeRabbit.

Business Actions:

- Implement Serge for code review to improve code quality.
- Adhere to data privacy policies when using Serge's codebase.

[Source](#)

2026-06-24

Mistral OCR 4 state-of-the-art

OCR Document Intelligence Enterprise AI RAG

Dev Actions:

- Mistral OCR 4 RAG pipeline document-heavy enterprise applications

- OCR 4 document digitization workflow
- Mistral OCR 4 Azure Document Intelligence

2026-06-23

Beyond Siri: Here are the practical AI features coming to your iPhone in iOS 27

📱 Siri: 🧠 AI 📱 iOS 27

TechCrunch 📱 AI 📱 Apple 📱 iOS 27 📱
📱 Siri
📱 Apple 📱 AI capabilities 📱
📱 iPhone 📱 on-device 📱
📱 Apple Intelligence 📱
📱 AI 📱
📱 Apple 📱 chatbot 📱
📱 AI 📱
📱 iPhone 📱 --- iOS 27 📱 AI 📱
📱 Apple 📱 iPhone 📱 Apple
Developer documentation 📱 AI API 📱 iOS 27 SDK 📱 on-device
model integration 📱 Apple Intelligence
📱 iOS 27 AI 📱 app strategy 📱 Apple
platform 📱 update 📱 Apple Intelligence
features 📱 iOS 27 release 📱: medium — 📱 Mobile AI,
On-device AI, Apple Intelligence

Mobile AI On-device AI Apple Intelligence iOS

Tech Impact: iOS 27 📱 AI 📱 Apple 📱 iPhone 📱

Dev Actions:

- 📱 Apple Developer documentation 📱 AI API 📱 iOS 27 SDK
- 📱 on-device model integration 📱 Apple Intelligence

Business Actions:

- 📱 iOS 27 AI 📱 app strategy 📱 Apple platform 📱
- 📱 update 📱 Apple Intelligence features 📱 iOS 27 release

Liquid AI Ships LFM2.5-230M with llama.cpp, MLX, vLLM, SGLang, and ONNX Support for On-Device Inference

Liquid AI [announces](#) LFM2.5-230M: [a](#) 230M [open-weight](#) On-Device Inference [model](#) llama.cpp, MLX, vLLM

Liquid AI [announces](#) LFM2.5-230M [a](#) 230 [million](#) open-weight [model](#) [designed](#) [for](#) [on-device](#) [inference](#) [across](#) [multiple](#) [frameworks](#) [including](#) [llama.cpp](#), [MLX](#), [vLLM](#), [SGLang](#), and [ONNX](#). The [model](#) [features](#) [a](#) LFM2 [architecture](#) [with](#) [a](#) LIV convolution 8 layer [and](#) [a](#) Grouped-Query Attention 6 layer [for](#) [inference](#) [on](#) [CPU](#) [and](#) [GPU](#). It [achieves](#) [a](#) 213 token [per](#) [second](#) [on](#) [Samsung](#) Galaxy S25 Ultra [and](#) [a](#) 42 token/s [on](#) [Raspberry](#) Pi 5. The [model](#) [was](#) [pre-trained](#) [on](#) [19](#) [billion](#) [tokens](#) [and](#) [post-trained](#) [with](#) [SFT](#) [and](#) [distillation](#) [on](#) [LFM2.5-350M](#), [DPO](#) [and](#) [multi-domain](#) [reinforcement](#) [learning](#) [framework](#). It [supports](#) [llama.cpp](#), [MLX](#), [vLLM](#), [SGLang](#) [and](#) [ONNX](#) [with](#) [a](#) memory footprint [of](#) [293-375](#) MB. The [model](#) [is](#) [designed](#) [for](#) [agentic](#) [AI](#) [on](#) [edge](#) [devices](#) [and](#) [mobile](#) --- LFM2.5-230M [is](#) [designed](#) [for](#) [agentic](#) [AI](#) [on](#) [edge](#) [devices](#) [and](#) [mobile](#) [for](#) [cloud](#) [API](#) [tool](#) [use](#) [and](#) [data](#) [extraction](#).

Edge AI

On-Device AI

LLM

Agents

Tech Impact: LFM2.5-230M [is](#) [designed](#) [for](#) [agentic](#) [AI](#) [on](#) [edge](#) [devices](#) [and](#) [mobile](#) [for](#) [cloud](#) [API](#) [tool](#) [use](#) [and](#) [data](#) [extraction](#)

Dev Actions:

- [Liquid](#) AI [announces](#) LFM2.5-230M [on](#) [Hugging](#) Face [with](#) [a](#) function calling [use](#) case [demo](#)
- [Unislo](#)th/TRL LoRA fine-tuning recipe [on](#) [Liquid](#) AI [for](#) [domain](#) [adaptation](#)

Business Actions:

- [Liquid](#) AI [API](#) [calls](#) [on](#) [on-device](#) LFM2.5-230M [for](#) [data](#) [extraction](#) [volume](#) [demo](#)
- [Liquid](#) AI [deployment](#) [on](#) [edge](#) [devices](#) [for](#) [IoT](#) [latency](#) [and](#) [protect](#) [data](#) [privacy](#)

[Source](#)

2026-06-28

8
'VIPA International Film Festival 2026'

Digital Media Content Creation EdTech

Tech Impact:

- Tech Partner

- Media Tech

- Source

Thumbsup

MEDIUM

2026 SME x Influencer EP 6

'2026' SME x Influencer EP 6

SME x Influencer EP 6 SME AI '2026' AI Influencer Marketing Authentic Connection Creator Audience Engagement AI Generated Content SME Social Proof Community Building AI Content --- AI-generated Content Authentic Brand Voice Competitive Advantage SME AI Tools Authenticity Influencer Content Fake Engagement Detection Platform SME Authentic Marketing Campaign SME Brand Story Community Influencer Micro-community Engaged Macro Influencer Followers Engagement

Digital Marketing

AI Content

SME

Tech Impact: AI-generated Content Authentic Brand Voice Competitive Advantage SME

Dev Actions:

- AI Tools Authenticity Influencer Content Fake Engagement Detection
- Platform SME Authentic Marketing Campaign

Business Actions:

- SME Brand Story Community
- Influencer Micro-community Engaged Macro Influencer Followers Engagement

[Source](#)

2026-06-28

□□□□□□ Ambani □□□□□□ AI □□□□□□ □□□□□ □□□□□□□□

AI Market India Tech Telecom AI Emerging Markets

Dev Actions:

- ### Business Actions:

- Source

2026-06-28

Nobel laureate John Jumper is leaving DeepMind for rival Anthropic

Anthropic CEO Dario Amodei says John Jumper is leaving DeepMind to join Anthropic

TechCrunch reports that John Jumper, a Nobel laureate and former DeepMind researcher, is leaving DeepMind to join Anthropic. Jumper, who worked at DeepMind for over a decade, is expected to join Anthropic in 2024. He will be leading Anthropic's AI research efforts, which are focused on developing more powerful and safe AI models. Jumper's departure from DeepMind is seen as a significant move, as he is one of the top researchers in the field. Anthropic is a leading AI research company, and Jumper's joining is expected to strengthen their research capabilities. Jumper's research at DeepMind focused on protein structure prediction, a field that has seen significant advances in recent years. His work on AlphaFold, a deep learning model for protein structure prediction, has been highly influential. Jumper's move to Anthropic is expected to be a key factor in Anthropic's efforts to develop more powerful and safe AI models. Anthropic is currently working on developing a new AI model, which is expected to be more powerful than GPT-4. Jumper's research at Anthropic is expected to focus on developing more powerful and safe AI models, which is a key goal for the company. Jumper's research at Anthropic is expected to be a key factor in Anthropic's efforts to develop more powerful and safe AI models. Anthropic is currently working on developing a new AI model, which is expected to be more powerful than GPT-4. Jumper's research at Anthropic is expected to focus on developing more powerful and safe AI models, which is a key goal for the company. Jumper's research at Anthropic is expected to be a key factor in Anthropic's efforts to develop more powerful and safe AI models. Anthropic is currently working on developing a new AI model, which is expected to be more powerful than GPT-4. Jumper's research at Anthropic is expected to focus on developing more powerful and safe AI models, which is a key goal for the company.

[AI Research](#)[Talent Acquisition](#)[AI for Science](#)[Biotech AI](#)

Tech Impact: Jumper's move from DeepMind to Anthropic is expected to strengthen Anthropic's AI research capabilities.

Dev Actions:

- Anthropic is expected to develop a new AI model, which is expected to be more powerful than GPT-4.
- Anthropic is expected to focus on developing more powerful and safe AI models, which is a key goal for the company.

Business Actions:

- Anthropic is expected to focus on developing more powerful and safe AI models, which is a key goal for the company.
- Anthropic is expected to focus on developing more powerful and safe AI models, which is a key goal for the company.

[Source](#)

2026-06-28

When the Trump administration cracks down on Anthropic, who benefits?

Who will benefit if Trump cracks down on Anthropic?

TechCrunch reports that the Trump administration is considering cracking down on Anthropic, a leading AI company. This move is part of a broader strategy to control AI technology, specifically targeting export controls. The administration is reportedly looking at Fable 5 and Mythos 5, which are AI models developed by Anthropic. The move is seen as a defensive security measure, similar to what OpenAI has faced. Anthropic's competitive advantage in the AI industry is at risk, and the move could lead to a fallback model strategy. The export control policy is expected to have a high impact on the AI industry, particularly in the areas of AI Policy, Export Control, and Geopolitics.

[AI Policy](#)[Export Control](#)[Geopolitics](#)[Cybersecurity](#)

Tech Impact: Export control risks Anthropic's competitive advantage in the AI industry.

Dev Actions:

- regulatory AI model model
- fallback model strategy

Business Actions:

- export control policy AI
-

[Source](#)

2026-06-29

AI 80% 4

AI 80% - 4

World Economic Forum (WEF) Summer Davos 2026 80% AI 4 AI Data Centers AI Liquid Cooling, Nuclear Power Data Centers Energy-efficient AI Models WEF AI Investment --- AI Data Center Investment Energy-efficient AI Models Model Compression, Quantization Pruning Green AI Frameworks Carbon-aware Computing AI AI Infrastructure Energy Policy AI Data Centers high — AI Infrastructure, Energy, Data Center

AI Infrastructure Energy Data Center WEF

Tech Impact: AI AI

Dev Actions:

- Energy-efficient AI Models Model Compression, Quantization Pruning
- Green AI Frameworks Carbon-aware Computing AI

Business Actions:

- AI Infrastructure
- Energy Policy AI Data Centers

 12,000 AI,
 10,000

AI Policy EdTech AI Research China

Tech Impact: [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] AI [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] AI Talent Pool [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] []

- ☐ Open Source AI Projects ☐
- ☐ AI ☐

- STEM AI
- Partnership Talent Pipeline AI

2026-06-23

2026-06-24

Data Center [NVIDIA 45]

Data Center [] [] NVIDIA [] [] 45 []

NVIDIA [] Data Center [] AI Factory [] Liquid Cooling [] 45 [] GPU [] 'Hot Water Cooling' [] [] [] AI [] Data Centers [] WEF [] 4 [] AI [] NVIDIA Rubin Architecture [] Platform [] Hot Water Cooling [] --- [] Hot Water Cooling [] NVIDIA [] Data Center [] AI Infrastructure [] NVIDIA Rubin Architecture [] Thermal Management APIs [] Hot Water Cooling [] Carbon Footprint [] AI Models [] Hardware [] Data Center [] Hot Water Cooling [] Default Design [] AI Infrastructure [] Energy Efficiency KPIs [] ROI Calculation

GPU Data Center AI Infrastructure NVIDIA

Tech Impact: [] Hot Water Cooling [] NVIDIA [] Data Center []

Dev Actions:

- [] NVIDIA Rubin Architecture [] Thermal Management APIs [] Hot Water Cooling
- [] Carbon Footprint [] AI Models [] Hardware []

Business Actions:

- Data Center [] Hot Water Cooling [] Default Design
- [] AI Infrastructure [] Energy Efficiency KPIs [] ROI Calculation

[Source](#)

2026-06-23

AI News Pipeline v2.0 | 2026-06-29 | 33 articles |
2026-06-29T03:55:26+07:00