

■ AI NEWS DAILY — 2026-07-26

■ 25 articles | AI Trends: 11 | Tech: 4 | Thailand: 4 | Global: 6 | Pipeline v2.0

■ TOP INSIGHTS TODAY

- NVIDIA ██████████ SK Hynix ██████████ HBM

■■■■■■■■■■ AI Factory ■■■■■■■■■■ 5 ■■■■■■■■■■











- **Prentis, new AI lab co-founded by Reid Hoffman, Mark Pincus in talks to raise \$100M**

[illegible]

- **Anthropic launches Opus 5**

[illegible]

- **AMD takes on Nvidia with its Helios AI rack-scale system**

Microsoft deploy Helios ██████████ real performance data ██████████
AMD claim 30% advantage — █ 2027 ██████████ NVIDIA moat ██████████

- **AegisAI, founded by former Google security execs, lands \$36M to stop AI-driven spear phishing**

Al spear phishing threat enterprise — attack AI personalization scale traditional email filter

■ AI TRENDS (11 articles)

[HIGH] [BREAKTHROUGH] [TechCrunch] Prentis, new AI lab co-founded by Reid Hoffman, Mark Pincus in talks to raise \$100M

Prentis ██████████ computer-use ██████ next AI mega-market ████████████████████ coding — █████ Hive-32B
 ████████████████████ 10x █████ frontier API ████████████████████ █████ disrupt enterprise software ████████████████████
 Prentis ████████████████████ 2026 █████ Ritankar Das (CEO, founder Titan holding co.), Reid Hoffman
 (LinkedIn co-founder, Greylock, ex-Microsoft board) █████ Mark Pincus (Zynga founder) ████████████████████ train
 ████████████████████ office workflow ████████████████████ ████████████████████ AI agent █████ control
 computer ██████████ ██████████ Hive-32B ████████████████████ OpenAI GPT-5.4 █████ Claude Opus 4.6 █████
 WindowsAgentArena █████ ScreenSpot-v2 benchmark ████████████████████ task ██████████ frontier API █████ 10
 ████████████████████ OpenAI, Google DeepMind, Meta, Tencent, Alibaba
 ████████████████████ \$50M ██████████ healthcare, manufacturer, clothing manufacturer

Impact: Prentis ■■■■ bet ■■■ computer-use AI ■■ outpace coding AI ■■■■■ market size — ■■■■■■■■■
■■■■■■■■■■ enterprise software market ■■■ white-collar

[HIGH] [TechCrunch] Anthropic launches Opus 5

Opus 5 ■■■■ strategic value play ■■■■ Anthropic — developer ■■■■■■■■■■ near-frontier capability ■■■■■■■■■■ Fable 5 ■■■■ Anthropic ■■■■■■■■ Claude Opus 5 ■■■■■■■■■■ Fable 5 (■■■■■ top tier ■■■■■■■■■■) ■■■■■■■■■■ \$5/\$25 ■■■■■■■■■■ token ■■■■■■■■■■ Opus 4.8 Opus 5 ■■■■■■■■■■ Opus 4.8 ■■■■ benchmark ■■■■ knowledge ■■■■ coding safety constraints ■■■■■■■■■■ Fable 5 ■■■■■■■■■■

Impact: Opus 5 ■■■■ cost barrier ■■■■■■■■ near-frontier AI — team ■■■■■■■■ Fable 5 ■■■■■■■■■■ Opus 5 ■■■■■■■■■■

Source: <https://techcrunch.com/2026/07/24/anthropic-launches-opus-5/>

[LOW] [TechCrunch] Anthropic updates Claude voice mode with more capable models

Voice AI ■■■■■■■■■■ battleground ■■■■■■■■■■ Anthropic ■■■■ OpenAI — ■■■■■■■■■■ voice mode ■■■■ timeframe ■■■■■■■■■■ Anthropic ■■■■■■■■ Claude voice mode ■■■■ upgrade ■■■■■■■■■■ power voice conversation ■■■■ powerful ■■■■ latency ■■■■■■■■■■ nuanced conversation ■■■■■■■■ Claude voice mode ■■■■■■■■■■ ChatGPT voice mode ■■■■ OpenAI ■■■■■■■■■■ desktop ■■■■

Impact: Voice AI ■■■■ interface ■■■■■■■■■■ AI ■■■■ adopt ■■■■ daily workflow ■■■■■■■■■■ — Anthropic ■■■■ OpenAI ■■■■ invest ■■■■ voice improvement ■■■■■■■■■■

Source: <https://techcrunch.com/2026/07/23/anthropic-updates-claude-voice-mode-with-more-capable-models/>

[HIGH] [ArXiv cs.AI] The Autonomous Agency Scale: A Behavioral Framework for Measuring Self-Directed Behavior in AI Systems

AAS ■■■■ missing piece ■■■■ AI safety toolbox — ■■■■■■■■■■ standardized way ■■■■■■■■ AI ■■■■■■■■ autonomy ■■■■■■■■■■ risk AAS ■■■■■■■■■■ behavioral rubric ■■■■■■■■■■ Samuel Presgraves ■■■■■■■■ Autonomous Agency Scale (AAS) ■■■■ behavioral framework ■■■■■■■■■■ self-directed behavior ■■■■■■■■ AI ■■■■■■■■ simple reactive system ■■■■■■■■ highly autonomous agent ■■■■■■■■■■ rubric ■■■■■■■■■■ dimension ■■■■ goal persistence, resource acquisition, deceptive capability, ■■■■ boundary testing AAS v0.2.1 ■■■■■■■■■■ GitHub ■■■■■■■■ assessment file ■■■■■■■■ OpenAI/Hugging Face breach event ■■■■■■■■■■ standardized way ■■■■ autonomy ■■■■ deploy production agent

Impact: AAS ■■■■ practical tool ■■■■ team ■■■■ build AI agent ■■■■■■■■ evaluate ■■■■ production deploy — ■■■■■■■■ agent ■■■■ autonomy ■■■■■■■■■■ risk profile ■■■■■■■■

Source: <https://arxiv.org/abs/2607.17947>

[HIGH] [ArXiv cs.AI] Can We Break LLMs Out of Self-Loops? Fine-Grained Reasoning Control with Activation Steering

Activation Steering ■■■■ technique ■■■■■■■■■■ developer ■■■■ fine-grained control ■■■■■■■■ LLM behavior ■■■■■■■■■■ retrain — ■■■■■■■■■■ practical ■■■■■■■■ agent self-loop problem ■■■■■■■■■■ UC San Diego, Adobe Research ■■■■ UNSW ■■■■■■■■ fine-grained reasoning control ■■■■■■■■ LLM ■■■■■■■■ activation steering ■■■■■■■■■■ LLM agent ■■■■■■■■■■ self-loop ■■■■■■■■■■ agentic workflow ■■■■■■■■ model ■■■■■■■■■■ terminate ■■■■ continue reasoning ■■■■■■■■■■ identify activation patterns ■■■■■■■■■■ self-loop behavior ■■■■ steer ■■■■■■■■ pattern ■■■■■■■■■■ inference ■■■■■■■■■■ fine-tune model ■■■■

Impact: Self-loop ■■■■■■■■■■ critical ■■■■ production AI agent ■■■■■■■■■■ consume token ■■■■■■■■■■ deliver result Activation Steering ■■■■ technique ■■■■ developer ■■■■■■■■ apply ■■■■■■■■■■...

[MED] [ArXiv cs.AI] Verify, Repair, Repeat, or Stop? Robust Stopping for Noisy Verify-Repair Loops in LLM Agents

Stopping criteria ██████████ reliability ███ cost ███ production AI agent — ██████████
 framework ███ practical ██████████ agentic AI ███████ LLM agent
 ███████ verify-repair loop ██████████ repair ██████████ failure
 ██████████ robust stopping mechanism ██████████ noisy verification signal
 ██████████ real-world scenario ███ code testing, document validation, ███████ API interaction
 ██████████ probabilistic stopping criteria ███ balance ██████████ false stop
 (████████████████████) ███ infinite loop (████████████████████) ██████████
 complement ██████████ self-loop paper ███ Sheldon Yu et al.

Impact: Stopping criteria █████ under-studied ██ AI agent design ████████████████████ token cost, reliability, ███ user experience ██████████ framework ███ implement █████

Source: <https://arxiv.org/abs/2607.17641>

[MED] [ArXiv cs.AI] Mechanistic Attention Guidance for Agent Memory Refinement

[illegible]

Impact: ■■■■■ long-running AI agent ■■■■ large memory store memory refinement ■■■ accurate ■■■ efficient
■■■■ critical factor ■■■■■■■■■ performance ■■■■■ degradation ■■■■■■■■■■■■■■■■■■■

Source: <https://arxiv.org/abs/2607.17621>

[MED] [ArXiv cs.AI] Exploratory and Assimilating Reflection: Reflective Recall Cycle for Long-term Memory

■■■■ reflect ■■■■ recall ■■■■ memory retrieval ■■■■ — pattern ■■■■ episodic
 memory ■■■■ context ■■■■ cue memory ■■■■ relevant Ganesh Senrayan et al. ■■■■ Reflective
 Recall Cycle (RRC) ■■■■ direct retrieval ■■■■ reflection step ■■■■ ■■■■ reflect ■■ current context
 ■■■■ generate cues ■■■■ meaningful ■■■■ query memory store ■■■■ retrieval
 ■■■■ raw query ■■■■ ■■ 2 mode: Exploratory Reflection ■■■■ connection
 ■■■■ ■■■■ Assimilating Reflection ■■■■ knowledge ■■■■ context
 ■■■■ ■■■■ long-horizon task ■■■■ AI companion scenario

Impact: RRC better memory architecture agent performance better base model
long-horizon task — agent session

Source: <https://arxiv.org/abs/2607.17879>

[MED] [BREAKTHROUGH] [ArXiv cs.AI] The Shared Discovery Paradox: How a One-Answer Rule Turns Better Information into Worse Search

Paradox ██████████ multi-agent system ██████████ — sharing information
 ██████████ coordination mechanism █████ balance exploration █████
 exploitation Yohei Nakajima (BabyAGI creator) ██████████ discover paradox █████ multi-agent
 search: ████████ system █████ one-answer rule (agent ████████████████████ share ████████████████████
 agent ████████████████████) ██ agent ████████████ search quality
 ████████████████████ ████████ early false positive █████ cascade effect ████████████████████ ████████████ code █████
 interactive guide █████ GitHub ████████████ █████ implication █████ design █████ multi-agent RAG, parallel agent search,
 █████ distributed knowledge systems

Source: <https://arxiv.org/abs/2607.18045>

[illegible]

 reshape  multi-agent AI —  SAP, Oracle, Microsoft

Dynamics disruption

Source: <https://arxiv.org/abs/2607.17331>

████████████████████ bridge ABM simulation (██████████ social science █████ policy simulation) █████ LLM
agent — ███████████████████ AI simulator ███████ complex social system █████████████████ Scuola
Normale Superiore █████ DTU █████████ traditional Agent-based Models ███████ rule-based agent
██████████████ behave ███ LLM-powered agent ███████ emergent behavior
██████████████████████████████████████ agentic ABM █████ result █████ realistic █████ █████████████████ unpredictability
████ LLM █████ model verification ███ statistical model checking approach
██████ verify property █████ agentic ABM ███

Source: <https://arxiv.org/abs/2607.17948>

Attie ■■■■ attempt ■■■■ Bluesky ■■■■■■■■■■ AI-native social intelligence layer ■■ open protocol — ■■■■■■■■
Twitter/X ■■■■■■ API Attie ■■■■■■■■■■ Bluesky ■■■■■■■■ Attie
■■■■■■■■■ AI chatbot ■■■■■■■■ ■■■■■■■■■■ open social research tool
■■■■■■■■■■■■■■■■■■■■ trending topics, ■■■■■■■■■■ sentiment, ■■■■ summarize conversation ■■
AT Protocol ■■■■ ■■ API ■■■■■■■■ developer ■■■■■■ application ■■ Attie ■■■■ ■■■■■■■■ differentiate ■■■■ X
■■■■■■■■■ API ■■■■■■■■■■■■■■■■■■■■■

Source: <https://techcrunch.com/2026/07/24/blueskys-ai-assistant-attie-expands-into-an-open-social-research-tool/>

[LOW] [TechCrunch] OpenAI's new voice mode makes it to the ChatGPT desktop app

Voice mode desktop OpenAI ChatGPT ambient AI assistant platform mobile OpenAI roll out voice mode mobile ChatGPT desktop app Windows macOS real-time desktop browser voice mode multimodal capability ChatGPT platform

Impact: Voice mode desktop use case hands-free AI interaction — feature developer knowledge worker

Source: <https://techcrunch.com/2026/07/24/openais-new-voice-mode-makes-it-to-the-chatgpt-desktop-app/>

[HIGH] [BREAKTHROUGH] [TechCrunch] AMD takes on Nvidia with its Helios AI rack-scale system

Helios most credible NVIDIA challenger rack-scale AI system — Microsoft adoption signal hyperscaler diversify NVIDIA dependency AMD Helios AI rack-scale integrate MI400 GPU, ROCm.ai software stack Infinity Fabric interconnect AMD Helios performance NVIDIA Vera Rubin flagship AI system NVIDIA 30% Microsoft launch customer deploy Helios rack hyperscale validation AMD hardware break NVIDIA's dominance

Impact: Microsoft deploy Helios real performance data AMD claim 30% advantage — 2027 NVIDIA moat

Source: <https://techcrunch.com/2026/07/23/amd-takes-on-nvidia-with-its-helios-ai-rack-scale-system/>

[MED] [BREAKTHROUGH] [TechCrunch] AI chip startup Etched defies skeptics, hits \$10.3B valuation from big-name investors

Etched bet Transformer architecture dominant — chip efficient GPU inference obsolete Etched ASIC chip optimize Transformer model throughput efficiency GPU generalist Transformer inference Etched valuation \$10.3B investor skeptic bet Transformer architecture architecture AI agentic paradigm flexible compute

Impact: Etched high-conviction bet Transformer longevity — Transformer dominant AI paradigm 5 Etched cost-efficient AI inference leader...

Source: <https://techcrunch.com/2026/07/23/ai-chip-startup-etched-defies-skeptics-hits-10-3b-valuation-from-big-name-investors/>

THAILAND AI & TECH (4 articles)

[HIGH] [Blognone] NVIDIA SK Hynix HBM AI Factory 5

NVIDIA-SK Hynix \$500B AI infrastructure supply chain HBM NVIDIA SK Hynix guaranteed revenue AI data center NVIDIA SK Hynix 5

(1) HBM (2) HBM (3) AI Factory 2 SK Telecom

NVIDIA DSX GPU Vera Rubin HBM SK Hynix 2026

Impact: NVIDIA supply chain HBM SK Hynix AI Factory 2GW refer...

Source: <https://www.blognone.com/node/151230>

[MED] [Blognone] Samsung Broadcom 2

Samsung-Broadcom \$200B Samsung Foundry — Broadcom 2nm capacity 5 custom AI chip demand Broadcom 2 5 2030 2nm Samsung Broadcom HBM 2.3D 2.5D AI Samsung Foundry TSMC

Impact: Samsung Foundry advanced node — Samsung TSMC custom AI chip Broadcom hyperscaler

Source: <https://www.blognone.com/node/151229>

[MED] [Blognone] Facebook Facebook Verified

Meta AI facial recognition trust layer platform — move 'verified badge' Facebook Facebook Verified (Meta Verified Facebook Verified Marketplace, Dating, Groups Facebook Feed

Impact: Meta AI facial recognition mass-market identity layer — biometric data collection privacy regulation PDPA

Source: <https://www.blognone.com/node/151232>

[LOW] [Blognone] AI

MLB 'AI governance in competitive contexts' — AI MLB (Major League Baseball) AI Dugout 1 3 AI tablet MLB MLB AI real-time AI

Impact: case study AI usage policy — 'AI'

Source: <https://www.blognone.com/node/151233>

[MED] [TechCrunch] Librarians are hosting viral 'Avoiding AI' workshops for people who are fed up with Big Tech

<https://techcrunch.com/2026/07/25/librarians-are-hosting-viral-avoiding-ai-workshops-for-people-who-are-fed-up-with-big-tech/>

<https://techcrunch.com/2026/07/24/as-us-weighs-response-to-chinese-ai-industry-urges-against-broad-open-weight-restrictions/>

Source: <https://techcrunch.com/2026/07/24/why-cognition-bought-poke-ai-personality-is-becoming-a-competitive-advantage/>

[MED] [TechCrunch] How AI guardrails are impeding the work of offensive cybersecurity researchers

■ tension ■ — AI ■ safe ■ defenders
■ attackers ■ uncensored model TechCrunch ■ offensive cybersecurity
researcher ■ AI tools ■ Claude Mythos ■
■ penetration testing, exploit development ■ vulnerability
analysis ■ professional context ■ researcher
■ open-source uncensored model ■
■

Impact: AI guardrail design ■ — ■ block too broad ■ legitimate security
professional ■ block too narrow ■ misuse

Source: <https://techcrunch.com/2026/07/23/how-ai-guardrails-are-impeding-the-work-of-offensive-cybersecurity-researchers/>

[HIGH] [TechCrunch] AegisAI, founded by former Google security execs, lands \$36M to stop AI-driven spear phishing

AI spear phishing ■ scale ■ human-crafted attack ■ AI ■ personalized
attack ■ — AegisAI ■ AI ■ threat landscape ■ AegisAI
■ Cy Khormae ■ Ryan Luo ■ Google security executive ■ \$36M
■ AI defense system ■ spear phishing ■ AI ■ AegisAI
■ email content, sender pattern ■ contextual signal ■ AI-generated phishing
■ traditional phishing ■ AI-powered
cyberattack ■ mainstream

Impact: AI spear phishing ■ threat ■ enterprise ■ — attack ■ AI ■
personalization ■ scale ■ traditional email filter ■

Source: <https://techcrunch.com/2026/07/23/aegisai-founded-by-former-google-security-execs-lands-36m-to-stop-ai-driven-spear-phishing/>

[MED] [TechCrunch] Google's Gemini nears billion-user milestone

Gemini ■ 1B users ■ AI assistant adoption ■ mainstream threshold —
■ early adopter ■ Google ■ earnings report ■ Gemini ■
milestone 1 ■ user ■ billion-user product ■ Search ■
Google Cloud ■ revenue growth +82% YoY ■ AI infrastructure demand
■ Google ■ monetize AI investment
■

Impact: Gemini ■ 1B users ■ Google AI ■ mainstream adoption ■ — ■ signal ■ AI
assistant ■ mass product ■ niche ■

Source: <https://techcrunch.com/2026/07/23/google-closes-in-on-another-billion-user-product-with-gemini/>