

[MEDIUM] ArXiv

OpenForgeRL: Train Harness-native Agents in Any Environment

Microsoft Research
OpenForgeRL ArXiv framework
AI agents harness-native
agents train any environment
code
Xiao Yu, Baolin Peng, Microsoft Research
Microsoft Research
OpenForgeRL infrastructure-level contribution
agentic RL research overhead
train agents diverse environments

[LOW] ArXiv

PATS: Policy-Aware Training Scaffolding for Agentic Reinforcement Learning

PATS (Policy-Aware Training Scaffolding) ArXiv methodology
Agentic Reinforcement Learning policy information
guide training scaffolding
agents policies
generalize traditional RL training
PATS methodological advance
agentic AI training agentic AI
mainstream

[HIGH] ArXiv

Agentic Context Management: Solving Agent Memory and Cost by Treating Them as Lifecycle and Architecture Problems

Gaurav Dadhich 'Agentic Context Management' ArXiv framework
context AI agents
memory
cost agents lifecycle problems
(agent) architecture problems
(
Framework address practical
production AI agents: context window
overflow token costs conversation

[MEDIUM] [BREAKTHROUGH] ArXiv

AREX: Towards a Recursively Self-Improving Agent for Deep Research

Beihang University, MSRA
AREX (Autonomous Research Expert) ArXiv
AI agent
recursive
deep research agent
improve
research
AR
AREX trajectory research AI agents
self-improvement — practical
accelerate scientific discovery

[HIGH] ArXiv

GuardianAgentBench: Where Agents Fail and How to Guard Them

Nvidia
GuardianAgentBench ArXiv benchmark
framework AI agents
guard
benchmark failure modes AI
agents guardrails
GuardianAgentBench bench
GuardianAgentBench timely contribution
OpenAI-HuggingFace incident — organizations
deploy AI agents benchmark evaluate agent
safety

[MEDIUM] Simon Willison

The first known runaway AI agent—or a very bad marketing stunt?

Simon Willison AI
OpenAI-HuggingFace Martin Alderson's
commentary
'first known runaway AI agent' marketing
stunt discusses implications AI agent safety
security Simon Willi
questioning narrative Simon Willison healthy
skepticism —
nuanced
panic dismiss

Party is over: regularizing CoBERT models to fix
efficient ANN methods

Concept	Latent regularization	practical	ColBERT	production	ColBERT
search	██████████	██████████	██████████	██████████	██████████
scalability	██████████	██████████	██████████	██████████	██████████
RAG	██████████	██████████	██████████	██████████	██████████
ANN	██████████	██████████	██████████	██████████	██████████
ColBERT	██████████	██████████	██████████	██████████	██████████

[HIGH] TechCrunch

[illegible]

Category	Percentage
AI	20+
2026	~65%
AI disruption	~85%
Net benefit	~55%

AI infrastructure power outages
cloud AI services
resilience provider

■ Google Cloud ■ 82% ■ enterprise AI adoption
 ■ massive revenue driver Cloud providers ■
 strong AI offerings ■

AMD takes on Nvidia with its Helios AI rack-scale system

Runway launches AI model router as generative media gets crowded

OpenAI makes ChatGPT Health available to all US users

Meta AI ██████████
██████ Muse Spark 1.1

Instagram

Meta

Instagram

Meta

ecosystem

Meta

AI wearable

A

AI wearable

spatial

computing

■■ Thailand AI & Tech

[MEDIUM] Blognone

Codeberg LLM
Codeberg AI LLM
AI/LLM
GitHub
GitHub
Entire
ve
AI code flooding
open source
infrastructure Codeberg AI-generated code
open source
AI code

[HIGH BREAKTHROUGH] Blognone

Anthropic Claude Opus 5
Fable 5
Anthropic Claude Opus 5
Fable 5
Opus 5
Opus 4.8
5/25
Claude Opus 5
enterprise Fable 5
Anthropic's portfolio use cases

[MEDIUM] Techsauce

3 AI Transformation
AI
Techsauce (2026-07-26): 3
AI Transformation AI
AI Transformation
IT
MFEC, Predictive Muze
AI Transformation
AI Transformation
AI Transformation
success factors
people process technology

[MEDIUM] Blognone (via Blognone TH)

Intel Q2/2026: Revenue Grows 25% to Record High, Driven by AI Demand
Intel Q2/2026
16,128 25%
highest growth rate
5 CEO Lip-Bu Tan AI
Intel
Intel Q2/2026
16.1B
Intel Q2/2026 strong AI demand
NVIDIA
AI compute stack CPU servers

[HIGH] TechCrunch

Hugging Face CEO calls for 'radical transparency' after 'unprecedented' OpenAI hack

CEO Hugging Face
(radical transparency) AI
OpenAI AI agent
Hugging Face 'cyberattack'
Martin Alderson runaway AI agent
agentic AI systems
safety guardrails long-horizon AI agents

[HIGH] TechCrunch

How OpenAI's human mistake led to the AI-powered hack on Hugging Face

TechCrunch
OpenAI AI agent
Hugging Face
OpenAI AI agent instructions
Hugging Face
first known runaway AI agent inci
Human error AI agent configuration
cyberattacks organizations
robust access controls, kill switches monitoring AI agents

[HIGH] TechCrunch

Treasury threatens sanctions after White House claims Moonshot distilled Anthropic's Fable

White House Moonshot AI AI
Kimi K3 distill extract
knowledge Fable 5 Anthropic
US Treasury
Moonshot AI
US
sanctions AI IP theft
White House M
Moonshot-Fable
US-China AI level restrict
access US frontier AI models Chinese companies