

Category	Count
AI Trends	5
Tech Trends	0
Thailand	0
Global	0
High Urgency	3
Breakthrough	5

TOP INSIGHTS TODAY

- Green SARC: Zero-Breach Budget Governance for Agentic AI**
First architectural framework proving 0% budget breach rate for agentic AI cost/carbon governance vs 91.5% for soft penalties; 47-55% token savings on 3,000 real plans
- SkillVetBench: LLM Security Vetting Catches What Static Scanners Miss 100%**
Zero false negatives on 78 malicious agent skills; static scanners miss 89-100% of prompt injection and memory poisoning threats in open-source LLM agent ecosystems
- ForgetEval: Mutation-Time LLM Hooks Essential for Production Agent Memory**
13-config architectural study shows control-plane placement determines which forgetting failures can be recovered; mutation-time hooks achieve 91.7% accuracy at \$0.17/385 cases

AI TRENDS (5)

DeepRoot: A KG-Coordinated Multi-Agent System for Therapeutic Reasoning over Historical Medical Texts

Source: ArXiv cs.AI | Research | 2026-06-14

DeepRoot multi-agent LLM ICML 2026 Knowledge Graph reasoning
drug discovery hallucination 87%
7-10% Shen Nong Ben Cao Jing
compound-disease pairs 10 21 R@20 (47.6%) 4.8% raw corpus LLM
grounding axis reasoning axis KG+LLM
Graph-only inference hallucinate rank
reasoning coherence ICML 2026 GenBio Workshop ACM CAIS
2026 biomedical AI multi-agent systems
AI
AI

Impact: KG-grounded multi-agent hallucination 10x drug discovery

Biz: Evaluate for pharma R&D; pipelines needing to extract leads from legacy documents | Track ICML 2026 GenBio track for follow-on work and partnership opportunities

Source: ArXiv cs.AI | Research | 2026-06-14

Impact: First provably zero-breach architectural governance framework for agentic AI cost and carbon budgets

Biz: Adopt Green SARC for compliance with emerging AI carbon reporting regulations | Use for FinOps enforcement in production agentic systems to cap runaway API costs

Source: ArXiv cs.AI | Research | 2026-06-14

Impact: Graph-based agentic RAG benchmark long multimodal documents

Biz: Evaluate for insurance, legal, and research document Q&A; requiring high accuracy | Benchmark against current PDF-RAG solutions to measure accuracy improvement potential

Source: ArXiv cs.AI | Research | 2026-06-14

SkillVetBench ██████████ SARS (Skill Agentic Risk Score) ████████████████████████████████ 5 ████████████████████████████████ LLM-as-Judge
 ████████████████████████████████ zero false negatives ████████ 78 agent skills ████████████████████████████████ zero false positives
 ████████ 22 skills ████████████████████████████████ static code scanners ████████████████████████████████ 100% ██████████
 instruction-layer threats ██████████ memory poisoning attacks ████████████████████████████████ 0 ████████ 9 ██████████
 ██████████ CVSS v4.0 ████████████████████████████████ leaderboard ████████████████████████████████ Hugging Face ████████████████████████████████ LLM
 evaluators ██████████ 35-95% ████████████████████████████████ ensemble scoring
 ████████████████████████████████ ecosystem ████████ open-source LLM agent skills ████████████████████████████████
 ████████████████████████████████ prompt

injection ■■■ memory poisoning

Impact: LLM-as-Judge eliminates false negatives that plague all static code-layer security scanners for agent skills

Dev: Integrate SkillVetBench evaluation before publishing agent skills to open-source marketplaces | Use SARS score as mandatory security gate in CI/CD pipelines for agentic AI applications

Biz: Audit all deployed open-source LLM agent skills against SARS criteria immediately | Establish mandatory security vetting policy for any third-party agent skills before production use

■ ■ Control-Plane Placement Shapes Forgetting: An Architectural Study of Agent Memory Across Thirteen System Configurations

Source: ArXiv cs.LG | Research | 2026-06-16

ForgetEval ██████████ LLM ███ agent memory pipeline ████████ forgetting failure modes
 ████████ recover ████████████████████ 13 system configurations ███ 385 adversarial test cases ████████
 mutation-time hooks ████████ 91.7-93.2% overall accuracy ██████████ deterministic-only approaches ██████████
 0% ██████████ canonicalization █████ intent-aware deletion ████████████ LLM ███ inscribe-time ████████ 100%
 canonicalization ██ intent-aware deletion ████████████ mutation-time hook ██████████
 78-85% ██████████ intent-aware deletion ████████████████████████████████ \$0.17 ███ 385 cases insight ██████████ production
 agent failures ████████████████████ forgetting failures █████ benchmark ██ recall ██████████
 ForgetEval (1000-case templated + 385 adversarial) ████████████████████ MIT license

Impact: First systematic study separating recall vs. control-plane benchmarking for agent memory systems

Dev: Add mutation-time LLM hook to agent memory pipelines to handle intent-aware deletion edge cases | Use ForgetEval benchmark suite to audit existing memory implementations before production deployment

Biz: Prioritize agent memory control-plane design in product roadmaps | Evaluate if current agentic products have adequate memory supersede/purge policies