

33	Total
15	AI Trends
5	Tech
11	Thailand
2	Global
13	High Urgency
15	Breakthroughs

TOP INSIGHTS — HIGH URGENCY TODAY

Anthropic is turning Claude Code's auto mode on by default

Auto Mode ██████████ Claude Code ████ signal ████████ autonomous AI coding agents
██████████ mainstream ██████████ developer workflow ██████████

Cloudflare launches Kitesurf, a browser built for AI agents

Kitesurf ██████████ AI agent infrastructure ██████████ market ██████████ ██████████ big players
██████████ Cloudflare ██████████

Jeff Dean and other top AI researchers are leaving Google to launch their own startup

Jeff Dean Google signal AI talent war frontier
AI startups Big Tech

OpenAI says it slowed Astra model development over security concerns

Entity	Count
OpenAI	10
Astra	9
self-governance	7
frontier AI labs	6
AI safety	5

Anthropic's Claude can now control your Mac, escalating the fight to build AI agents that actually do work

Claude Mac control █ Cowork+Code ████████ agentic ██████████
 ███ escalates ████████ OpenAI Operator

15 articles

[MEDIUM] TechCrunch | ■■■■■■■■■■ AI

The AI safety test is becoming a safety risk

AI

TechCrunch AI (AI safety evaluations) AI companies overfit safety benchmarks benchmarks red teamers pattern AI labs patch root cause evaluation methodology robust

Impact: AI safety evaluation methodology credibility safety claims frontier AI labs

AI Safety | Red Teaming | Benchmarks | AI Evaluation

[HIGH] [BREAKTHROUGH] VentureBeat | AI Agent

Anthropic's Claude can now control your Mac, escalating the fight to build AI agents that actually do work

Claude Anthropic Mac — AI Agent mainstream

Anthropic computer use Mac Claude Claude Cowork Claude Code Claude GUI Mac research preview Dispatch Claude Code output escalates AI agent OpenAI Operator Google Gemini computer use

Impact: Claude Mac control Cowork+Code agentic escalates OpenAI Operator

Anthropic | Claude | Mac Control | AI Agent

[HIGH] [BREAKTHROUGH] VentureBeat | AI

Google's new TurboQuant algorithm speeds up AI memory 8x, cutting costs by 50% or more

Google TurboQuant algorithm AI memory 8 50%

Google TurboQuant quantization AI model memory 8 inference 50% TurboQuant (weights) LLM quantization AI inference costs AI company models Google TurboQuant Google Cloud AI services open-source

Impact: TurboQuant cost of AI inference ROI AI adoption industry

Google | TurboQuant | Quantization | AI Memory

[HIGH] [BREAKTHROUGH] VentureBeat | AI

Mistral AI just released a text-to-speech model it says beats ElevenLabs — and it's giving away the weights for free

Mistral AI Text-to-Speech ElevenLabs weights

Topic	Frequency (approximate)
naturalness	30
ElevenLabs	10
voice AI	15
Mistral	25
model weights	15
open-source	25
ElevenLabs	10
OpenAI	5
voice AI market	25
Google	5
Cloud Chirp, IBM watsonx, OpenAI Speech	5
ElevenLabs	5
Mistral	15
open-weights strategy	30

Mistral AI | TTS | Text-to-Speech | ElevenLabs

AgentOPSD recursive self-distillation technique agentic AI on-policy reinforcement learning external reward signals human annotations AI agents iterative distill better-performing trajectories baselines agentic tasks tool use, web navigation coding code GitHub AgentOPSD step AI agents self-improve

Agentic RL | Self-Distillation | Recursive | LLM

paradigm RAG systems top-K retrieval
black-box operation interpretable agentic operations AI
retrieval strategy exact match, semantic similarity,
multi-hop reasoning documents RAG systems
enterprise explainability audit trails AI

RAG | Agentic Operations | Information Retrieval | Interpretable AI

DASH: ■■■■■ self-distillation ■■ on-policy ■■■■■ supervision horizon ■■■■■ reasoning models

DASH technique train reasoning models on-policy self-distillation supervision horizons dynamic divergence model outputs model supervisions diverge current policy training DASH adapt supervision horizon distribution model reasoning model train DASH self-distillation benchmarks self-improving AI systems

Impact: DASH contribution self-improving reasoning models frontier challenges AI

Reasoning Models | Self-Distillation | On-Policy | LLM Training

[MEDIUM] [BREAKTHROUGH] ArXiv cs.AI | AI

EnvACE: Internalizing Environment Dynamics via World Rehearsal for Agentic Reinforcement Learning

EnvACE: dynamics World Rehearsal Agentic Reinforcement Learning

EnvACE Huawei Noah's Ark Lab Shanghai Jiao Tong University agentic reinforcement learning agent internalize environment dynamics 'world rehearsal' real environment interactions agent generalize environments real-world experience EnvACE agentic AI agents adapt scenario

Impact: EnvACE's world rehearsal approach train agentic AI dependence expensive real-world interactions

Agentic RL | World Model | Reinforcement Learning | Environment Dynamics

[MEDIUM] ArXiv cs.AI | AI

FinEvo-Bench: A Longitudinal Benchmark for Self-Evolving Agents in Professional Financial Workflows

FinEvo-Bench: Benchmark self-evolving AI agents professional finance

Qwen DianJin Alibaba Cloud Beihang University FinEvo-Bench benchmark self-evolving AI agents professional financial workflows longitudinal market conditions financial agents static benchmarks FinEvo-Bench agent adapt evolve financial knowledge FinTech banking sectors deploy AI agents investment research, risk assessment financial advisory

Impact: FinEvo-Bench AI agents finance production deployment FinTech

Financial AI | Self-Evolving Agents | Benchmark | FinTech

[MEDIUM] ArXiv cs.AI | AI

From Siloed Algorithms to Compliance-First Agentic Platforms: A Multi-Layered Architecture for Hospital AI Systems

AI multi-layer compliance

self-evolving LLM agents 'skill contamination' agent catastrophic forgetting fine-tune neural networks pre-commit gating mechanism contaminate commit pre-commit gating overall agents safety self-evolving AI systems

Impact: Skill contamination safety concern self-evolving AI agents deploy production

Self-Evolving Agents | Skill Contamination | LLM Agent | Pre-Commit Gating

[MEDIUM] ArXiv cs.AI | AI

TRAJDEBUG: Tracing Error Lifecycle to Identify Critical Failures in Long-Horizon Agent Trajectories

TRAJDEBUG: long-horizon agents

Tsinghua University TRAJDEBUG framework debug AI agent trajectories long-horizon tasks AI agents TRAJDEBUG trace 'error lifecycle' trajectory root cause cascade failures agentic AI deploy production debugging capability reliability trust

Impact: TRAJDEBUG debugging deploy agentic AI production systems

Agentic AI | Debugging | Long-horizon | Agent Trajectory

[MEDIUM] ArXiv cs.AI | AI

TS-RAG: Retrieval Augmented Generation for Time Series Forecasting

TS-RAG: RAG Time Series LLM

TS-RAG Retrieval-Augmented Generation (RAG) time series forecasting LLMs distribution shift context historical patterns TS-RAG retrieve time series patterns database in-context examples LLM TS-RAG forecasting benchmarks finance, energy, supply chain time series forecasting

Impact: TS-RAG LLMs time series forecasting enterprise use cases

RAG | Time Series | Forecasting | LLM

TECH TRENDS

5 articles

[HIGH] [BREAKTHROUGH] TechCrunch | AI

Anthropic is turning Claude Code's auto mode on by default

Anthropic Auto Mode Claude Code — AI coding agent

[HIGH] [BREAKTHROUGH] TechCrunch | AI

[illegible]

Cloudflare | Kitesurf | Browser | AI Agents

[MEDIUM] TechCrunch | ■■■■■■■■■■ AI

Meta ■■■■■■ Muse Code: AI Agent ■■■■■■ codebase ■■■■■■■■■■■■■■ enterprise

Meta | Muse Code | AI Coding | Large Codebases

[HIGH] [BREAKTHROUGH] TechCrunch | ■■■■■■■■■■ AI

[illegible]

AI 4

Huawei AI+ Healthcare Summit Huawei Thailand Digital & AI Summit 2026
4 AI
AI computing
AI
AI for Health, Health for All
AI

Impact: AI Healthcare Huawei health tech AI
Big Tech

Huawei | AI Healthcare | Thailand | Smart Hospital

[MEDIUM] [BREAKTHROUGH] Blognone | AI

AMD Taalas Inference

AMD Taalas AI
Inference Taalas 219
Taalas AI
inference
Llama 3.1 Meta AMD
AI inference capability NVIDIA GPU AI
AI inference

Impact: Taalas AMD inference chips NVIDIA
AI workloads

AMD | Taalas | AI Chip | Inference

[MEDIUM] [BREAKTHROUGH] Blognone | AI

Anthropic Claude

Anthropic
Claude AI Google Amazon
AI
Claude
Anthropic GPU
NVIDIA inference
OpenAI Google DeepMind
Anthropic

Impact: Anthropic silicon stack
inference Claude

Anthropic | Chip | AI | Hardware

[MEDIUM] Blognone | AI

ChatGPT 1

OpenAI ChatGPT (weekly active users) 1 3 2022 Google Maps 5 OpenAI 1 2026 AI assistant OpenAI Impact: ChatGPT 1 AI assistant mainstream tool ChatGPT | OpenAI | 1 Billion Users

[HIGH] Blognone | AI Claude Code Anthropic Claude Code multi-agent workflows cross-session messaging Claude Code parallel development agent Impact: Cross-session messaging Claude Code multi-agent coding workflows Claude Code | Anthropic | AI | Cross-session

[HIGH] Blognone | AI DeepSeek API DeepSeek AI OpenAI API DeepSeek strategy (loss-leader strategy) AI LLM developer DeepSeek API Impact: DeepSeek API era LLM AI development DeepSeek | LLM | API Pricing

[HIGH] [BREAKTHROUGH] Blognone | AI OpenAI Astra Critical Safeguard

[illegible][illegible]

Volta cloud computing 1 Anthropic AI cloud infrastructure ecosystem neocloud GPU NVIDIA (Vera Rubin) Bitdeer inference Claude Volta AI demand compute infrastructure AI

Impact: Volta-Anthropic \$10B AI cloud infrastructure

Volta | Anthropic | AI Cloud | Neocloud

GLOBAL AI & TECH 2 articles

[HIGH] [BREAKTHROUGH] TechCrunch | AI

Jeff Dean and other top AI researchers are leaving Google to launch their own startup

Jeff Dean AI top Google startup

Jeff Dean Google Brain AI Google AI senior Google AI startup Google AI Jeff Dean TensorFlow, TPU chips Large Language Models AI Dean AI researchers top Big Tech startups upside

Impact: Jeff Dean Google signal AI talent war frontier AI startups Big Tech

Google | Jeff Dean | AI Startup | Talent

[MEDIUM] TechCrunch | AI

OpenAI acquires presentation startup NextSlide

OpenAI NextSlide presentation AI productivity tools Microsoft PowerPoint Google Slides OpenAI Work build suite of AI-powered productivity tools Microsoft 365 Google Workspace OpenAI NextSlide ChatGPT Work Microsoft OpenAI

Impact: NextSlide OpenAI Work strategy AI productivity suite Microsoft 365

OpenAI | NextSlide | Acquisition | Presentation AI