

Al
hyperscaler

Meet 'Code-as-World': An Agentic Loop That Rewrites Real Videos Into Executable MuJoCo Physics Programs [BREAKTHROUGH]

Impact: The world model is a key component of the system, enabling it to learn about the environment and predict future states.

Impact: frontier Continual Learning open-weight

本表展示了 Qwen3-72B 模型在评估意识（eval-awareness）和思维链（chain-of-thought）方面的表现。表格列出了不同提示词（prompt）下的得分，以及模型在不同任务上的表现。

Prompt	Score	Model
eval-awareness	0.98	Qwen3-72B
chain-of-thought	0.98	Qwen3-72B
capabilities-flavored	0.98	Qwen3-72B
safety-flavored	0.98	Qwen3-72B
Qwen3-32B	0.98	Qwen3-32B
FORTRESS capabilities-framing	0.98	Qwen3-72B
safety-framing	0.98	Qwen3-72B
CoT-prefill	0.98	Qwen3-72B
eval-awareness	0.98	Qwen3-72B

Tags: AI Safety, Eval-awareness, Interpretability, LLM

Impact: Evolution Strategies (ES) post-training LLM eval-awareness GRPO reasoning sequential GRPO-ES Pass@1 GRPO entropy collapse catastrophic forgetting ES post-training GRPO

Understanding Evolution Strategies for LLM Reasoning: Broader Reasoning Coverage than GRPO

Source: ArXiv ML (cs.LG) | Urgency: MEDIUM | 2026-08-27T23:48:39+07:00

Evolution Strategies (ES) post-training LLM GRPO reasoning sequential GRPO-ES Pass@1 GRPO entropy collapse catastrophic forgetting ES post-training GRPO

Tags: Evolution Strategies, GRPO, Reasoning, Post-training

Impact: Evolution Strategies post-training reasoning coverage LLM GRPO

Performance Foundations of Parallel & Distributed Reasoning Language Models

Source: ArXiv ML (cs.LG) | Urgency: MEDIUM | 2026-08-27T19:33:47+07:00

ETH Zurich Torsten Hoefler Maciej Besta Reasoning Language Models RLM DeepSeek-R1 o3 Kimi k1.5 Reinforcement Learning with Verifiable Rewards GPU-hours algorithm PPO GRPO compute parallelism data tensor pipeline sequence context expert parallelism disaggregated placement stage fusion asynchronous execution work-depth model RLM infrastructure AI

Tags: Reasoning Models, Distributed Systems, RLVR, GPU

Impact: infrastructure pipeline reasoning model

Puro-2B: Poor Lab's Qwen2-1.5B Trained on RTX 5090 within \$5090 [BREAKTHROUGH]

Source: ArXiv NLP (cs.CL) | Urgency: LOW | 2026-08-28T00:07:12+07:00

Puro-2B pretraining open-source LLM Llama-3.2-3B 1.5 SmolLM3-3B 7 RTX 5090 FP8 1.4 6,900 Qwen2.5-1.5B Puro Cost Scaling Law 4,400 5,090 Qwen2-1.5B

hyperball optimization curriculum model averaging Apache 2.0 Hugging Face

Tags: LLM Pretraining, Open-source, GPU, Cost Efficiency

Impact: LLM

TTPO: Test-Time Policy Optimization

Source: ArXiv NLP (cs.CL) | Urgency: MEDIUM | 2026-08-28T00:58:10+07:00

TTPO Test-Time Policy Optimization LLM ground-truth label majority-vote pseudo-label label TTPO On-Policy Self-Distillation Grouped RL label TTPO label 5 Qwen3-1.7B 38.0 45.2 test-time training 25.2 36.4 thinking

Tags: Test-Time Training, Reasoning, Self-Distillation, LLM

Impact: LLM label reasoning

LeVJEPa: Efficient & Scalable Video Pretraining without the Heuristics [BREAKTHROUGH]

Source: ArXiv Vision (cs.CV) | Urgency: MEDIUM | 2026-08-28T00:26:24+07:00

Yann LeCun LeVJEPa representation collapse target encoder exponential-moving-average encoder projector SIGReg collapse epoch LeVJEPa V-JEPa 2 compute 5.6 20.8 FLOPs LeVJEPa baseline 7.6 ImageNet-1K block-causal attention motion DINOv2 pretrain

Tags: Video Pretraining, JEPa, Self-Supervised Learning, Computer Vision

Impact: compute

UrbanGround: From Local Perception to Spatial Agency in a Real-Scale City

Source: ArXiv Vision (cs.CV) | Urgency: MEDIUM | 2026-08-28T00:59:33+07:00

UrbanGround 3 MLLM agent closed-loop agent agent

UrbanGround AI agent

Tags: MLLM Agents, Spatial Reasoning, Navigation, Benchmark

Impact: MLLM agent agent

TECH TRENDS (5)

Salesforce Anthropic Salesforce Claude

Source: Blognone | Urgency: MEDIUM | 2026-08-29T18:35:00+07:00

Salesforce Anthropic Claude
Claudeforce 29 2569
Salesforce in Claude
Claude
Cowork skill 37 Claude
Salesforce Claude inside Salesforce Claude
Agentforce reasoning Claude
AI Slack Slackbot Slack Code Claude Tag
Salesforce AI
Claude CRM Salesforce Alforce harness Headless 360
Data 360 Tableau Slack Anthropic
OpenAI Google AI
Salesforce Agentforce Claude
Salesforce in Claude

Tags: Salesforce, Anthropic, Claude, Enterprise AI

Impact: Claudeforce Claude AI CRM Salesforce Anthropic

OpenAI Cursor SpaceX Elon Musk

Source: Blognone | Urgency: HIGH | 2026-08-29T13:39:00+07:00

OpenAI Cursor AI
Cursor SpaceX
OpenAI OpenAI
Cursor SpaceX
Cursor OpenAI
Cursor OpenAI
12 2569
Cursor OpenAI
Composer Grok xAI
SpaceX Anthropic Compute
Claude Cursor OpenAI
AI

Impact: Gemini Omni AI

Building Custom Batched Ensemble Weather Forecasting with NVIDIA Earth2Studio

Source: MarkTechPost | Urgency: LOW | 2026-08-30T01:57:26+07:00

MarkTechPost 環境構築の要らないAIプラットフォーム (ensemble)
 NVIDIA Earth2Studio
 PyTorch CUDA Colab
 FCN GFS
 10
 API Earth2Studio iterator
 Zarr
 GFS RMSE
 fair CRPS
 AI

Tags: NVIDIA, Earth2Studio, Weather Forecasting, Tutorial

[illegible]

THAILAND (3)

Workday AI- HR

Source: Bangkok Biz AI | Urgency: MEDIUM | 2026-08-29T15:30:00+07:00

2023 年 10 月 10 日，Workday 宣佈將收購 Enterprise AI 平台開發商 Sana，這將使 Workday 成為第一家擁有完整 HR 技術堆棧的企業。Sana 的收購將為 Workday 提供強大的 AI 驅動力，使其能夠為企業提供更具針對性的 HR 服務。Workday 表示，Sana 的收購將為其客戶提供更具針對性的 HR 服務，並使其能夠為企業提供更具針對性的 HR 服務。Workday 表示，Sana 的收購將為其客戶提供更具針對性的 HR 服務，並使其能夠為企業提供更具針對性的 HR 服務。

Tags: Workday, Fusion Solution, HR Tech, Enterprise AI, Thailand

Impact: AI-powered misinformation can spread rapidly, leading to public confusion and harm. AI-generated content can be used to create fake news, deepfakes, and other deceptive information. AI can also be used to target individuals with personalized disinformation campaigns.

GLOBAL (4)

[illegible]

Source: Blognone | Urgency: HIGH | 2026-08-29T15:06:00+07:00

The Information 2022 AI Diffusion Rule NVIDIA Moonshot AI Kimi 3 Alibaba AI GPU

Tags: GPU Export, AI Diffusion Rule, China, Thailand, Trump

Impact: **AI Diffusion Rule** **NVIDIA** **GPU**

👉 **1** **LinkedIn** **AI Slop**

Source: Bangkok Biz AI | Urgency: MEDIUM | 2026-08-30T07:30:00+07:00

[illegible]

