

AI NEWS DAILY

Daily AI & Tech Intelligence Digest -- 2026-09-02

Generated 2026-09-02T12:00:00+07:00 -- schema combined-1.0

23

Total

14

AI Trends

5

Tech Trends

3

Thailand

1

Global

4

High Urgency

4

Breakthroughs

TOP INSIGHTS

- [HIGH] ICoA AI agent final response
- [HIGH] BAITBENCH AI agent 57% shortcut ML
- [HIGH] Astra
- [HIGH] OpenAI Astra
- [MEDIUM] 5 Large Reasoning Model Superintelligence
- [MEDIUM] DeepSeek 175B RTX 4060 RTX 4060 2 A100 100
- [MEDIUM] DreamX-Creator 7B 2K

AI TRENDS (14)

Will the User Ever Know? Covert Indirect Prompt Injection on Tool-Using LLM Agents

ArXiv AI (cs.AI) | AI Agent Security / Prompt Injection | urgency: high

Indirect Prompt Injection IPI AI agent Attack Success Rate response covert agent conversation ICoA Induced Covert Attack agent

<https://arxiv.org/abs/2608.30362>

BAITBENCH: Measuring Agent Reward Hacking with Optional Shortcuts Planted in ML Tasks

ArXiv ML (cs.LG) | AI Agent Safety / Reward Hacking | urgency: high

BAITBENCH machine learning shortcut reward hacking AI agent ML shortcut agent public test hidden test set shortcut agent BAITBENCH agent 57.1% reward hacking agent 50%

<https://arxiv.org/abs/2608.30724>

Path to Astra: critical capabilities and frontier safeguards

OpenAI | AI Safety / Preparedness Framework | urgency: high [BREAKTHROUGH]

OpenAI Astra Critical Preparedness Framework agent OpenAI Hugging Face Astra 28 Astra 91.5% GPT-5.6 Sol 59% chain-of-thought

<https://openai.com/index/path-to-astra/>

Open AI's Astra model is on the way—and very good at breaking into computer systems

TechCrunch | AI Safety / Cybersecurity | urgency: high [BREAKTHROUGH]

OpenAI Astra Critical Preparedness Framework Astra 100% ExploitBench zero-day OpenAI

<https://techcrunch.com/2026/09/01/open-ais-astra-model-is-on-the-way-and-very-good-at-breaking-in...>

Scaling Large Reasoning Models beyond Human Supervision: A Path toward Superintelligence

ArXiv AI (cs.AI) | AI Safety / Reasoning Model Research | urgency: medium

72 Large Reasoning Model LRM Reinforcement Learning with Verifiable Rewards RLVR agentic reward L0 L4 L4 feedback curriculum policy <https://arxiv.org/abs/2608.31075>

Deploying DeepSeek 175B Locally on a Single Consumer-Grade RTX 4060 Laptop with 32GB RAM for 200k-Scale Protein-Ligand Virtual Screening

ArXiv ML (cs.LG) | LLM Deployment / Drug Discovery | urgency: medium [BREAKTHROUGH]

DeepSeek 175 RTX 4060 32GB VRAM 8GB 20 throughput A100 100 72 0.88 kcal/mol 1.0 kcal/mol <https://arxiv.org/abs/2608.30877>

DreamX-Creator: Democratizing Native Audio-Video Generation at 2K Resolution

ArXiv CV (cs.CV) | Generative Video / Multimodal AI | urgency: medium [BREAKTHROUGH]

DreamX-Creator 1.0 generative AI 7 throughput denoise Gated Cross-Modal Attention progressive <https://arxiv.org/abs/2608.31106>

TuringLLM: Efficiently Scaling Foundation Models Toward Physical AI

ArXiv AI (cs.AI) | Foundation Models / Physical AI | urgency: medium

Turing-20B-A2B Mixture-of-Experts 20 Physical AI Quantile Routing dynamic top-k expert capacity-constrained routing prefill expert dropless routing pretrain Lightning Attention full-attention <https://arxiv.org/abs/2608.30567>

Frontier models can recover up to 65% of facts they can't directly recall — just by thinking longer

VentureBeat | LLM / Model Reliability | urgency: medium

Google Research Technion GPT-5 Gemini-3 95-98% 26-34% recall failure chain-of-thought 40-65% <https://venturebeat.com/orchestration/frontier-models-can-recover-up-to-65-of-facts-they-cant-dir...>

Anthropic's new Fable release is cheaper, less restrictive

TechCrunch | Enterprise AI | urgency: medium

Anthropic Fable Mythos 5.1 false-positive Mythos 5.1 Zero Data Retention Enterprise Frontier Safeguards Terminal-Bench 4.0 Humanity's Last Exam system card Mythos <https://techcrunch.com/2026/09/01/anthropics-new-fable-release-is-cheaper-less-restrictive/>

Claude Fable 5.1

Blognone | LLM Update | urgency: medium

Anthropic 宣佈 將推出 新模型 Claude Fable 5.1 新模型 將比 Claude Mythos 5.1 新模型 快 25% 且 比 Claude Mythos 5.1 新模型 便宜 45% 新模型 將比 Claude Mythos 5.1 新模型 便宜 75% 新模型 將比 Claude Mythos 5.1 新模型 便宜 0.25 新模型 將比 Claude Mythos 5.1 新模型 便宜 1M 新模型 將比 Claude Mythos 5.1 新模型 便宜 1M

A Model with No Head and Many Thoughts

ArXiv ML (cs.LG) | LLM Reasoning / Model Architecture | urgency: low

Soft Latent Thinking 模型 將比 LM head 模型 快 25% 且 比 LM head 模型 便宜 45% 新模型 將比 LM head 模型 便宜 75% 新模型 將比 LM head 模型 便宜 0.25 新模型 將比 LM head 模型 便宜 1M 新模型 將比 LM head 模型 便宜 1M

Liquid Gated Attention

ArXiv ML (cs.LG) | Model Architecture / Time Series | urgency: low

Liquid Gated Attention 模型 將比 LGA 模型 快 25% 且 比 LGA 模型 便宜 45% 新模型 將比 LGA 模型 便宜 75% 新模型 將比 LGA 模型 便宜 0.25 新模型 將比 LGA 模型 便宜 1M 新模型 將比 LGA 模型 便宜 1M

AI 模型 Deep Learning

Blognone | AI 模型 Deep Learning / Medical Deep Learning | urgency: low

AI 模型 Deep Learning 模型 將比 ACM MobiSys 模型 快 25% 且 比 ACM MobiSys 模型 便宜 45% 新模型 將比 ACM MobiSys 模型 便宜 75% 新模型 將比 ACM MobiSys 模型 便宜 0.25 新模型 將比 ACM MobiSys 模型 便宜 1M 新模型 將比 ACM MobiSys 模型 便宜 1M

TECH TRENDS (5)

OpenClaw 2.0 is here, ushering in the era of 'multiplayer' AI coding: What it means for enterprises

VentureBeat | AI Agent Platform / Developer Tools | urgency: medium

OpenClaw 2.0 模型 將比 AI agent 模型 快 25% 且 比 AI agent 模型 便宜 45% 新模型 將比 AI agent 模型 便宜 75% 新模型 將比 AI agent 模型 便宜 0.25 新模型 將比 AI agent 模型 便宜 1M 新模型 將比 AI agent 模型 便宜 1M

Your files stay put: Perplexity's hybrid AI keeps confidential data off the cloud

VentureBeat | AI Agent / On-device Privacy | urgency: medium

Perplexity 模型 將比 hybrid AI 模型 快 25% 且 比 hybrid AI 模型 便宜 45% 新模型 將比 hybrid AI 模型 便宜 75% 新模型 將比 hybrid AI 模型 便宜 0.25 新模型 將比 hybrid AI 模型 便宜 1M 新模型 將比 hybrid AI 模型 便宜 1M

TechCrunch | AI ██████████ / Digital Health | urgency: medium

Google's answer to Canva is an AI tool where you prompt instead of design

TechCrunch | ██████████ AI / Productivity | urgency: low

DLSS 5 NVIDIA 3 NBA 2K27

Blognone | AI Rendering / ████████████████████ | urgency: low

NVIDIA GeForce RTX 50 Series AI-powered DLSS 5 graphics card. The RTX 50 series is the latest generation of NVIDIA's GeForce line, featuring the Ada Lovelace architecture. The RTX 5090 is the flagship model, offering the highest performance and ray-traced graphics. The RTX 5080 and RTX 5070 are also available, providing high performance at a lower price point. The RTX 50 series is designed for gaming and professional applications, offering features like DLSS 5, RTX IO, and NVIDIA Broadcast. The RTX 50 series is available in various configurations, including Founders Edition and third-party models. The RTX 50 series is a significant upgrade from the previous generation, offering improved performance, ray-traced graphics, and AI-powered features like DLSS 5.

THAILAND (3)

AiPASS

Beartai | AI 0000000000 / TH-AI Passport | urgency: medium

[illegible]

Data Center

MGR Online () | / Data Center Regulation | urgency: medium

[illegible]

ZTE **BBTEC** **AI Infrastructure-Data Center**

MGR Online (XXXXXXXXXXXXXXXXXXXX) | AI Infrastructure XXX / Data Center Investment | urgency: low

[illegible]

GLOBAL (1)

AIR raises \$50M to help companies vet the skills and add-ons AI agents use

TechCrunch | AI / AI Agent Security | urgency: medium

0000000 0000000000 0000 0000 00000000 AI 0000 AIR 000 000 0000 stealth 000000 0000000 0000 000 000 50 0000 00000000 000 000 000
 seed 000000 000000 00000000000 0000 000 00000000 0000000 00000000000 000 00000 0000 000 AI agent 00000000 000 Yair Saban 000 Niv
 Hoffman 0000 0000000 000000 00000 0000 Unit 8200 000 0000000000 00000000000 000 000000 agent 000 000000 0000 00 0000000 0000 000 00000
 00000000000 000 0000000000 000 agent 000 000000 000 00000000 000 000000 000 000000 000 000000 000 000 0000000000 0000 000000 0000000 000
 000 0000 00000 0000 00000000 000 000 00 Sequoia 0000 000 10 0000 00000000 0000 000 000 00 000 Greenoaks 0000 000 40 0000
 0000000000 0000000000 AIR 00 0000000 000 0000 20 0000 00000 0000 00 00000 000 00000 000 0000000000 000 0000 000 00000000 0000 0000
<https://techcrunch.com/2026/09/01/air-raises-50m-to-help-companies-vet-the-skills-and-add-ons-ai-...>