AI NEWS DAILY

Daily AI & Tech Intelligence Digest — 2026-10-19

Generated 2026-10-19T10:30:00+07:00 — schema vcombined-1.0

Total	AI Trends	Tech Trends	Thailand	Global	High Urgency	Breakthroughs
13	5	2	2	4	5	4

TOP INSIGHTS

[HIGH] Microsoft AI □□□□□□□□□□□□□□□□ MAI-Base-2 □□□□□□□□□□□□□□□□ □□□□□□□□□□□□□□□□□□□□□□□□□

Microsoft is building an end-to-end frontier pre-training stack, reducing dependence on partners and adding competition at the model layer.

[HIGH] Jack Clark □□□□ Anthropic □□□□□□□ IMF-World Bank □□□□□□□□□□□ □□□□□□□□□□□□□□□□□□□□□ AI □□□□□□□□□□□□□□□□□

A frontier-lab co-founder framing AI progress as near a self-accelerating phase at a major policy venue in Bangkok raises the urgency of national AI strategy.

[MEDIUM] Terminal-RSI: □□□

Verifier-driven recursive self-rewriting gives a recipe for improving mid-size open models as terminal and coding agents.

[MEDIUM] MLExpert: □□

A 35B MoE agent nearly doubles medal rate on MLE-Bench Lite, suggesting small specialised agents can rival larger frontier models in automated ML work.

[HIGH] OpenCode Console □□□□□□ □□□□□□□□□ API key □□□□□□□□ □□□□□□□□□□□□□□□□□□□□□□□

Centralised API-key sharing in dev tooling is a high-value target; stolen LLM keys can be monetised and used for distillation.

[HIGH] □□□□□□ Volker Rath Field CISO □□□ Cloudflare □□□□□ Agent □□□□□□ □□□□□□□□□□□□□?

AI agent accountability and Shadow AI are shifting from tech problems to governance, policy-as-code and zero-trust enforcement; relevant to Thai regulatory moves.

AI TRENDS (5)

Terminal-RSI: Recursive Self-Improving Terminal Agents via Recursive Self-Rewrite and Policy-Refine [BREAKTHROUGH]

Source: ArXiv AI (cs.AI) | Urgency: MEDIUM | 2026-10-15

□□□□□□□□□□□□□□□□□ Terminal-RSI □□□ □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ Recursive Self-Rewrite □□ □□□□ Recursive Policy-Refine □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ critic □□□□□□□□□□□□□□□□□□□□□□□□□ advantage □□□□□ TB-GAE, □□□□□□□□ Family Dense Reward (FDR) □□□□ Success-Prioritized Reward (SPR) □□ □□□□□□□□□□□□□□□□□□□□□□□□□□□□□ Qwen3.8-27B □□□□□□□□□□□□□□□□□□ □□□□□□□□□□□ Terminal-RSI-27B □□□□□□□□□□□□□□□□□ Terminal-Bench 4.0 □□□□ 19.84%, Terminal-Bench 2.1 □□□□ 86.52%, Terminal-Bench 3.0 □□□□ 16.95% □□□□ Terminal-Bench Hard □□□□ 79.00% □□ LHTB □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ 0.21 □□□□□ 0.37 □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ 0/46 □□□□□ 5/46 □□□□□ SWR100 □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ 3.00% □□□□□ 6.00% □□□

Tags: arXiv, terminal agents, Terminal-RSI, RL, Qwen3.8

Impact: Verifier-driven recursive self-rewriting gives a recipe for improving mid-size open models as terminal and coding agents.

LEASH: Uncertainty-Aware Weak-to-Strong Escalation and Oversight for Computer-Use Agents

Source: ArXiv AI (cs.AI) | Urgency: MEDIUM | 2026-10-15

□□□□□□□□□□ LEASH □□□□□□□□□□□□□□□□ training-free □□ □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ frontier □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ open-weight □□ □□ □□□□□□□□□□□ shadow observer □□□ API □□□□ black-box □□□□□□□□□□□□□□□□□□□ □□□□□□□□ □□ □□□□ shadow observer □□□ hidden states □□□ □□□□□□□ activation □□□□□□□□□□□□□□□□□□□□□□□□□ forward pass □□□□□□□ □□ OSWorld □□□□□□□□□□□□□□□□□□□□□ LEASH □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ 6.55 □□□□ 11.94 □□□□□□ □□□□ Qwen3.5-35B-A3B □□□□□□□□□□ 73.66% □□□□□□□□□□ 0.4117 □□□□□□□□□□□□□□□□□ □□□□□□□□ Claude Opus 5 □□□□□□□□□□□□□□□ (72.90%) □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ 19.2%

Tags: arXiv, LEASH, computer-use agents, OSWorld, uncertainty

Impact: Uncertainty-based routing can match frontier computer-use agents at roughly 19% lower inference cost, with built-in oversight.

MLExpert: Towards a Generalist Agentic Model for Autonomous Machine Learning Engineering [BREAKTHROUGH]

Source: ArXiv AI (cs.AI) | Urgency: MEDIUM | 2026-10-15

□□ (MLE) □□□□□□□□□□□□□□□□□□□ AI □□□□□□□□□□ AI □□□ □□□□□□□□□□□□□□□□ MLExpert □□□□□□□□□□□□□□□□□□□□□□□□□□□□ Mixture-of-Experts □□□□□ 35B □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ □□□□ Kaggle-MLE Factory □□ harness □□□□□□□□□□□□□□□□□□□□□□□□□□□□□, Hierarchical agentic SFT □□ MLE, Windowed Interaction RL (WIRL) □□□ □□ MLE-Bench Lite □□ 37.9% □□□□□ 72.7% □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□ Kimi-K3 [illegible] GLM-5.3 [illegible] NatureBench Lite [illegible] agent scaffold [illegible]

Tags: arXiv, MLExpert, MLE-Bench, MoE, agentic RL

Impact: A 35B MoE agent nearly doubles medal rate on MLE-Bench Lite, suggesting small specialised agents can rival larger frontier models in automated ML work.

Realtime-Bench: Can Frontier Models Solve Real-time Interactive Problems?

Source: ArXiv AI (cs.AI) | Urgency: MEDIUM | 2026-10-15

[illegible] Realtime-Bench [illegible] Wall Clock [illegible] Paused Clock [illegible] Token Clock [illegible] Wall Clock [illegible] 44% [illegible] Pass^3 [illegible] 36% [illegible]

Tags: arXiv, Realtime-Bench, real-time agents, benchmark, latency

Impact: Real-time interaction remains a clear weakness of frontier models (best 44% success), exposing a gap for robotics, gaming and live-agent use cases.

Microsoft AI unveils MAI-Base-2 base model, far ahead of its first generation [BREAKTHROUGH]

Source: Blognone | Urgency: HIGH | 2026-10-18

Microsoft AI [illegible] MAI-Base-2 [illegible] pre-training [illegible] post-training [illegible] Mixture-of-Experts [illegible] 2.8T [illegible] 80B [illegible] MAI-Base-1 [illegible] 1T [illegible] 35B [illegible] Microsoft [illegible] GPU GB200 [illegible] 16,000 [illegible] 38 [illegible] MAI-Base-1 [illegible] DeepSeek-V4-Pro-Base [illegible] open-weight [illegible] post-training [illegible] Kimi K3, Qwen3.8-2.4T-A95B [illegible] MiMo-V2.6-Pro-RL [illegible] MAI-Base-2 [illegible] Microsoft [illegible]

Tags: Microsoft AI, MAI-Base-2, MoE, pre-training, GB200

Impact: Microsoft is building an end-to-end frontier pre-training stack, reducing dependence on partners and adding competition at the model layer.

TECH TRENDS (2)

OpenCode Console hacked, attackers extract API keys and run up large charges

Source: Blognone | Urgency: HIGH | 2026-10-18

OpenCode [illegible] OpenCode Console [illegible] [illegible] [illegible] API key [illegible] [illegible] API [illegible] [illegible] OpenCode [illegible] [illegible] OpenCode Go [illegible] API key [illegible] [illegible] Console [illegible] API key [illegible] OpenCode [illegible] API key [illegible] [illegible] distillation [illegible] 1.4% [illegible] [illegible] [illegible] OpenCode [illegible] Anthropic, OpenAI [illegible] Google [illegible] [illegible] API key [illegible]

Tags: OpenCode, API key, LLM, data breach, distillation

Impact: Centralised API-key sharing in dev tooling is a high-value target; stolen LLM keys can be monetised and used for distillation.

Google Meet adds post-meeting screen with Gemini summaries to email and Google Tasks

Source: Blognone | Urgency: LOW | 2026-10-19

Google [illegible] Google Meet [illegible] [illegible] Gemini [illegible] [illegible] [illegible] [illegible] (action items) [illegible] Google Tasks [illegible] Gemini [illegible] [illegible] Google Tasks [illegible] Google Workspace [illegible] [illegible] Android [illegible] [illegible] [illegible] [illegible] iOS [illegible] [illegible] Android [illegible] [illegible] [illegible] [illegible] Google [illegible] Gemini [illegible] [illegible] Workspace

Tags: Google Meet, Gemini, Google Tasks, Workspace

Impact: Embedding Gemini into post-meeting workflow reduces manual follow-up work and deepens Workspace lock-in.

THAILAND (2)

Anthropic's Jack Clark tells IMF-World Bank in Bangkok we have under a decade to prepare for AI progress [BREAKTHROUGH]

Source: Techsauce | Urgency: HIGH | 2026-10-18

Jack Clark [illegible] Anthropic [illegible] Per Jacobsson Foundation Lecture [illegible] IMF-World Bank 2026 [illegible] [illegible] Inside the Event Horizon [illegible] [illegible] Anthropic [illegible] 2 [illegible] 2026 [illegible] 2025 [illegible] 9.3 [illegible] [illegible] AI [illegible] [illegible] Navier-Stokes [illegible] Erdős [illegible] Humanity's Last Exam [illegible] Claude Opus 5.5 [illegible] 67.7% [illegible] 8% [illegible] [illegible] FrontierMath [illegible] 94% [illegible] GPT-6 Astra [illegible] Anthropic [illegible] GDP [illegible] 44 [illegible] 2030 [illegible] AI [illegible] [illegible] AI [illegible] [illegible] AI [illegible]

Tags: Jack Clark, Anthropic, IMF-World Bank, Bangkok, Per Jacobsson Lecture

Impact: A frontier-lab co-founder framing AI progress as near a self-accelerating phase at a major policy venue in Bangkok raises the urgency of national AI strategy.

Cloudflare Field CISO Volker Rath: when an AI agent makes a mistake, who is accountable?

Source: Techsauce | Urgency: HIGH | 2026-10-19

Techsauce □□□□□□□□ Volker Rath Field CISO □□□ Cloudflare □□□□□ GovWare 2026 □□□□□□□□□□□ □□□ AI □□□□□□□□□□□□□□□□□□□□ □□□ Post-Mythos □□□□□□□□□□□□□□□□□□□□ □□□□ □□□□□□□□□□□□□□□□□□□□□□□□ □□□□□□□ Zero Trust □□□□□□□□□ NIST □□ 2020 □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□□□□□□□□□□□□□□□ WAF □□□□□□□ □□□□□□□□□□ Cloudflare Security Signals □□□□□□□□□□ 98% □□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□□ Shadow AI □□□□ Cloudflare □□□□□□□ Zero Trust □□□□□□□ AI Gateway □□□□□□□□□□□□□□□□□□□□□□□□□□ 10% □□□□□□□□□□□□□□□□□ □□□□□□□□ Policy-as-Code □□□□□□□□□□□□ PDF □□□□□□□□□□□□□□□□□□□□□□□□□□□ agent □□□□□□□□□□□□□□□□□□□□□□□ □□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□□ CISO □□□ AI □□□□□□□□□□□□□□□□□□□□□□ □□□□□□□□□□□□□ AI □□□□□□□□□□□□□□□□□□□□□□□□□□□

Tags: Cloudflare, Zero Trust, AI agents, Shadow AI, GovWare 2026

Impact: AI agent accountability and Shadow AI are shifting from tech problems to governance, policy-as-code and zero-trust enforcement; relevant to Thai regulatory moves.

GLOBAL (4)

[Rumor] Accenture starts capping employees' Claude usage at $29 per month

Source: Blognone | Urgency: MEDIUM | 2026-10-18

Blognone □□□□□□□□□□□□□□□□□□□□ 404 Media □□□ Accenture □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□□□□□□□□□□□□ Claude □□□ Anthropic □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ 29 □□□□□□□□□□□□□□□ □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ 1,000 □□□□□□□□□□□□□□□ 404 Media □□□□□□□□□□□□□□□□□□□□□□□□□□□ Claude Enterprise Usage □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□ 29 □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ Accenture □□□□□□□□□□□□□□□□□□□□□ □□□ Work Breakdown Structure (WBS) □□□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ □ □□□□□□□□□□□□□□□□□ AI □□□□□□□□□□ □□□ Blognone □□□□□□□ Microsoft □□□ Meta □□□□□□□□□□□□□□□□□□ Claude □□□□□□□ □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□

Tags: Accenture, Claude, Anthropic, AI cost, rumor

Impact: Enterprise AI spend is shifting from unlimited experimentation to per-seat budgets, favouring in-house or cheaper models.

AI makes hacking banks easier: IMF panel warns global finance has 12 months to respond

Source: Techsauce | Urgency: HIGH | 2026-10-18

Techsauce □□□□□□□□□□□□□□ IMF □□□□□□ AI □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□ □□□□□□□□□□□□ Anthropic, FINMA □□□□□□□□□□□□□□□□□□, □□□□□□□□□□□□□□□□ □□□ BIS □□□□□□□□□□□□□□□□ AI □□□□□□□□□□□□□□□□□□□□□ □□ □□□□□□ Anthropic □□□□□□□ Claude □□□□□□□□□□□□□□□□ 50% □□□□□□□□□□□□□□□□□□ 80-90% □□□□□□□□□□□

[illegible] Glasswing [illegible] 130,000 [illegible] 30,000 [illegible] FINMA [illegible] MOVEit [illegible] BIS [illegible] Pix [illegible] 12 [illegible]

Tags: IMF, BIS, FINMA, Anthropic, Glasswing, cyber risk

Impact: AI-accelerated offence compresses attack timelines; systemic third-party concentration makes financial-sector resilience urgent.

Anthropic commits 5% of its shares to a new Public Benefit Fund

Source: Anthropic News | Urgency: MEDIUM | 2026-10-18

Anthropic [illegible] Anthropic Public Benefit Fund [illegible] wholly-owned LLC [illegible] 5% [illegible] 2% [illegible] 50% [illegible] 80% [illegible] 20% [illegible] AI [illegible] Anthropic [illegible] 200 [illegible] AI [illegible] Dario Amodei [illegible]

Tags: Anthropic, Public Benefit Fund, equity, governance

Impact: Frontier labs are experimenting with equity-based public-benefit structures as trust in AI declines and IPO scrutiny rises.

Disney agrees to hire the majority of Character.AI's workforce

Source: TechCrunch AI | Urgency: MEDIUM | 2026-10-18

TechCrunch [illegible] Disney [illegible] Character.AI [illegible] 60 [illegible] 60 [illegible] Bloomberg [illegible] “[illegible]” [illegible] Karandeep Anand [illegible] Character.AI [illegible] CTO [illegible] Disney [illegible] Character.AI [illegible] Disney [illegible] 20 [illegible] Disney [illegible] stalking horse bid [illegible] Character.AI [illegible] 14 [illegible]

Tags: Disney, Character.AI, bankruptcy, acqui-hire

Impact: Consumer companion-chatbot economics plus child-safety liability can sink a well-funded AI startup, with talent absorbed by incumbents.

AI NEWS PUBLISHER v2.1 — Generated 2026-10-19T10:30:00+07:00 — 13 articles processed